

Long-Term-Care Utility and Late-in-Life Saving

John Ameriks

The Vanguard Group

Joseph Briggs

Federal Reserve Board of Governors

Andrew Caplin

New York University and National Bureau of Economic Research

Matthew D. Shapiro

University of Michigan and National Bureau of Economic Research

Christopher Tonetti

Stanford Graduate School of Business and National Bureau of Economic Research

Older wealth holders spend down assets much more slowly than predicted by classic life-cycle models. This paper introduces health-dependent utility into a model with incomplete markets in which preferences for bequests, expenditures when in need of long-term care, and ordinary consumption combine with health and longevity uncertainty

This research is supported by a program project grant from the National Institute on Aging, P01-AG026571. The Vanguard Group supported the data collection through the VRI. Vanguard's Client Insight Group and Ipsos were responsible for implementing the survey and provided substantial input into its design. The design of the VRI benefited from the collaboration and assistance of Wandí Bruine de Bruin, Alycia Chin, Brooke Helpie McFall, Minjoon Lee, Mi Luo, and Ann Rodgers, as part of the program project, from Annette Bonner of The Vanguard Group, and from the research staff at Ipsos. The views expressed herein are

Electronically published April 23, 2020

[*Journal of Political Economy*, 2020, vol. 128, no. 6]

© 2020 by The University of Chicago. All rights reserved. 0022-3808/2020/12806-0008\$10.00

to explain saving behavior. To sharply identify motives, it develops strategic survey questions (SSQs) that elicit stated preferences. The model is estimated using these SSQs and wealth data from the Vanguard Research Initiative. The desire to self-insure against long-term-care risk explains a substantial fraction of the wealth holding of many older Americans.

I. Introduction

The elementary life-cycle model predicts a strong pattern of dissaving in retirement. Yet this strong dissaving is not observed empirically.¹ Establishing what is wrong with the simple model is vital for understanding behavior and the optimal design of Social Security, Medicare, Medicaid, retirement savings plans, and private insurance products. Given the aging of the US population, identifying the determinants of late-in-life saving behavior is an increasingly important endeavor.

At present, there is no consensus on why there is so little spend-down of assets. Explanations typically involve longevity risk and either bequest motives, precautionary motives associated with high late-in-life health and long-term-care (LTC) expenses, or both, but the quantitative contribution of each motive is in debate. LTC needs are a quantitatively plausible driver of saving because the chance of needing care is nontrivial and the costs of care are large. Around one-third to one-half of older Americans will spend time in a nursing home, and, according to Genworth (2016, also available at <http://www.longtermcare.gov>), one year in a private room in a nursing home averages \$92,000 and ranges from \$55,000 to \$250,000.² This range in costs partially reflects large variation in quality of care and comfort, which suggests that spending when in need of LTC involves a choice with an intensive margin reflecting the utility of spending when in need of care.

Both bequest and health-related motives can by themselves generate the high levels of observed savings late in life. Estimates of the importance of these motives range widely. Kopecky and Koreshkova (2014) find LTC expenses to be a significant driver of savings; De Nardi, French, and Jones (2010) find medical expenses to be important in replicating

those of the authors and do not necessarily reflect the views of The Vanguard Group or the Federal Reserve Board of Governors. For documentation of the VRI, including a dynamic link to the survey instrument, see <http://ebp-projects.isr.umich.edu/VRI/>.

¹ Soto, Penner, and Smith (2009) find that the wealthiest 20% of the HRS (Health and Retirement Study) sample report rising net worth until age 85, and Love, Palumbo, and Smith (2009) and Poterba, Venti, and Wise (2013) similarly show that household wealth is relatively stable at later ages, absent divorce or death of spouse.

² See Brown and Finkelstein (2008), Friedberg et al. (2014), and Hurd, Michaud, and Rohwedder (2017) for the probability of nursing home use.

the slow spend-down of wealth. Others, including Hubbard, Skinner, and Zeldes (1994) and Palumbo (1999), estimate the contribution of such expenses to late-in-life savings to be low. Bequests as a saving motive have been studied extensively, with Kotlikoff and Summers (1981) and Hurd (1989) providing early analysis of the effect of a bequest motive on wealth decumulation. Most recent empirical work models an end-of-life bequest motive with the nonhomothetic utility functional form proposed by De Nardi (2004). While such studies broadly agree that the bequest motive is present and active primarily for richer individuals (and even found by Kopczuk and Lupton 2007 to be present for individuals without children), its quantitative importance is debated. Lockwood (2018) estimates a bequest utility function with high marginal utility in the relevant domain, which can by itself largely explain the high saving rates of the elderly. Others, such as De Nardi, French, and Jones (2010) and Ameriks et al. (2011), estimate the motive to be somewhat weaker, diminishing more rapidly in the bequest level, leaving more room for health-related saving motives to drive saving among the sample who could feasibly purchase private LTC. See De Nardi, French, and Jones (2016b) for a survey of the literature on savings after retirement.

We provide new estimates of the relative importance of bequest and precautionary motives. First, we find that there is a significantly stronger preference for spending when in need of LTC than for leaving a bequest. By estimating utility functions, we can show that the marginal utility from an additional dollar spent is higher when in need of LTC than when leaving a bequest at all levels of expenditure. Also, reported spending when faced with a hypothetical direct trade-off between leaving a bequest and spending on self when in need of LTC reveals a strong desire to spend when in need of LTC, even at the direct cost of leaving a smaller bequest. Second, we show that features of the health and economic environment lead both of these motives to contribute substantially to saving over the life cycle. Simulations of the full model, which takes into account not just preferences but also the probabilities of health and death events and the availability of government-provided means-tested care, show that LTC-related and bequest motives are roughly coequal in determining late-in-life saving, although the effect of each saving motive on saving behavior varies by wealth.

Our results derive from four interrelated innovations. The first concerns the modeling strategy. We build a heterogeneous-agent incomplete-markets model of individuals, who save precautionarily when faced with health risks, the potential need for LTC, and an uncertain life span. People value consuming, leaving a bequest, and receiving LTC if they need it. From at least as early as Arrow (1974), economists have postulated that utility may be state dependent and that health may be an important state that determines utility. A critical element of our modeling strategy is to

allow for an intensive margin of LTC expenditure that is valued using a non-homothetic LTC-state-dependent utility function. Specifically, we model LTC utility symmetrically with the bequest utility function proposed by De Nardi (2004). Existing models are asymmetric in this regard, allowing bequests to have a flexible state-dependent utility yet treating LTC either as a fixed expense or as a portion of standard consumption, sometimes with a distinct marginal utility multiplier. Allowing this additional flexibility reflects the distinctive nature of the spending options and desires when in need of LTC. We also model the option for individuals to utilize the publicly provided insurance against LTC and health risks. We incorporate these social insurance programs as means-tested consumption floors, with a separate provision for LTC and non-LTC health states. While clearly social insurance programs provide consumption for the US population with no wealth, the perceived value of these social insurance programs affects savings across the wealth distribution.³

Our second innovation is one of measurement. We develop a series of strategic survey questions (SSQs) to help identify preference parameters (see Barsky et al. 1997 and Ameriks et al. 2011). Our use of this variant of the stated-preference method is related to work by van der Klaauw and Wolpin (2008) and van der Klaauw (2012), by the use of nonbehavioral data to estimate structural model parameters. In contrast, those papers use subjective-expectations data, while we implement SSQs that elicit stated strategies in structured hypothetical scenarios to estimate utility functions.

Our third innovation is our estimation approach. We estimate a structural life-cycle model in the spirit of Gourinchas and Parker (2002), De Nardi, French, and Jones (2010), French and Jones (2011), and Lockwood (2018). Most importantly, we use not only standard behavioral data but also nonstandard SSQ data to jointly estimate risk aversion, LTC utility parameters, and bequest utility parameters.⁴ While our favored specification leverages both types of data and estimates the model jointly combining wealth and SSQ moments, we also provide separate estimates

³ For evidence on the effect of such social insurance programs, see, e.g., Hubbard, Skinner, and Zeldes (1995), Scholz, Seshadri, and Khitatrakun (2006), Brown and Finkelstein (2008), Ameriks et al. (2011), and Braun, Kopecky, and Koreshkova (2017).

⁴ In previous literature, two primary empirical strategies have been used to identify health-state-dependent utility. The first is to use panel data to analyze health profiles over time and the corresponding levels of consumption (Lillard and Weiss 1997) or utility proxies (Finkelstein, Luttmer, and Notowidigdo 2013). The primary alternative has been to use a compensating-differentials approach (Viscusi and Evans 1990; Evans and Viscusi 1991), asking survey respondents how much they would need to be paid to compensate for hypothetical health risks, often in the context of physically dangerous jobs. Finkelstein, Luttmer, and Notowidigdo (2009) provide an overview of the empirical strategies used to identify preferences in poor health states. See Hong, Pijoan-Mas, and Ríos-Rull (2015) for an alternative method using Euler equations to estimate the effect of health on the marginal utility of consumption.

using moments of the wealth distribution alone and SSQs alone.⁵ Additionally, in order to perform this estimation, we develop a method to efficiently compute optimal policies that builds on the endogenous-grid method of Fella (2014). The individual's value function is non-concave in wealth, which generates discontinuous optimal saving policies. Efficient computation of the optimal policies using this modified endogenous-grid algorithm and parallelization greatly facilitates estimation of the model.

Our final innovation relates to the sample. We derive our results in the context of a new sample, the Vanguard Research Initiative (VRI), that explicitly targets the half of older Americans with nontrivial financial assets and obtains high-quality measurements of wealth. While this sample is not randomly selected from the US population, we document in detail, in Ameriks et al. (2015), that it has much in common with the appropriately conditioned HRS. In section IX.B, we show that our estimated model predicts well saving patterns out of sample in the HRS population. Analysis of the VRI enables the use of SSQs while simultaneously providing new data on a previously undersampled relevant population for the question at hand.⁶

Ultimately, we examine the implications of the estimated preferences for savings and expenditure profiles. These estimates suggest that spending when in need of LTC is highly valued on the margin and contributes substantially to savings.⁷ It is striking that this broad conclusion holds not only for the estimates based on both wealth and SSQ data but also for

⁵ Our estimation procedure is related to other strategies in which stated choices are used to estimate models in the same manner as are data on observed choices. For a recent example that highlights the similarities and differences between the classic stated-preference and our strategic survey methodologies, see Blass, Lach, and Manski (2010). The closest paper to ours in this dimension is by Brown, Goda, and McGarry (2016), who use a related survey methodology to study the degree to which there exists health-state-dependent utility. As in our paper, they do find evidence of state dependence. They do not estimate a state-dependent utility function.

⁶ Ameriks et al. (2018c) uses the VRI to study the demand for insurance against the risk of needing help with the activities of daily living (ADLs). It examines both model-implied and stated demands for idealized LTC insurance (LTCL) and compares them to demand for LTCL that the market provides. The findings point to substantial unmet demand for good LTCL and therefore point to imperfections and frictions in the insurance market. The current paper has no insurance market, so the hedge against risk is asset accumulation. Additionally, the current paper presents a method-of-moments estimator targeting wealth and SSQ data to estimate homogeneous preference parameters, while Ameriks et al. (2018c) targets SSQs exclusively to estimate different parameters for each respondent.

⁷ In contrast to our findings, Koijen, Van Nieuwerburgh, and Yogo (2016), who allow for similar motives but use different estimation methods and data on the observed demand for insurance products, find a strong preference for leaving a bequest and a lower marginal utility when in poor health. Similarly, while not featuring health-specific utility, Lockwood (2018) matches moments of the cross-sectional wealth distribution and LTCL demand by cohort and finds a strong preference to leave a bequest. The findings of low utility when in need of LTC result from matching low holdings of LTCL. Low LTCL ownership could

either type of data taken in isolation. With the flexible health-state utility functional form, even estimates targeting exclusively the traditionally used wealth moments suggest the importance of LTC for savings, but imprecisely. Combining SSQ and wealth moments confirms that much of late-in-life saving is driven by LTC-related desires by providing much sharper identification of preferences.

The remainder of the paper is organized as follows. Section II develops the model; section III describes the financial and demographic data in the VRI; section IV details the SSQs; section V describes the estimation methodology that allows us to estimate the structural life-cycle model without (and with) data on observed behavior; section VI presents our baseline parameter estimates, obtained by matching both wealth and SSQ moments; section VII examines the resulting behavioral implications of the estimated preferences; section VIII compares our baseline estimates to those obtained by exclusively targeting either wealth moments or SSQ moments to disentangle the relative contribution of the SSQs and also compares the baseline estimates to those found in the literature; section IX performs sensitivity analysis by reestimating the model under alternative model and parameter assumptions; section X concludes.

II. The Model

A. Individual States

Individuals are heterogeneous over wealth ($a \in [0, \infty)$), age ($t \in \{55, 56, \dots, 108\}$), income-age profile ($y \in \{y_1, y_2, \dots, y_5\}$, with $y_k = \{y_k(t)\}_{t=55}^T$), sex ($g \in \{m, f\}$), health status ($s \in \{0, 1, 2, 3\}$), and health cost ($h \sim H_g(t, s)$). Time is discrete, and the life-cycle horizon is finite. Each period, consumers choose consumption (c), savings (d), and whether to use government care ($G \in \{0, 1\}$). The model groups consumers into five income groups with deterministic age-income profiles. For tractability, we abstract from labor-income risk, which, while essential to model for younger individuals, is less a determinant of behavior for older individuals out of the labor force or near retirement, who have a larger ratio of financial to human wealth.⁸ Assets receive a risk-free rate of return of $1 + r$ on savings. Thus, the only uncertainty an individual has is over health and death.

arise because the poor-quality LTCI products available in the market are worse than the idealized state-contingent assets in the model. See Ameriks et al. (2018c) for analysis of LTCI demand that makes the distinction between idealized and actual insurance products, which reconciles why individuals may have a strong preference to spend when in need of help with LTC and simultaneously not buy LTCI.

⁸ The model abstracts from labor supply decisions, including retirement. These labor market decisions are taken into account through the exogenous income profiles.

B. Health and Death

There are four health states: $s = 0$ represents good health, $s = 1$ represents poor health, $s = 2$ represents the need for LTC, and $s = 3$ represents death. People are defined to need LTC ($s = 2$) if they need help with ADLs (activities of daily living), such as bathing, eating, dressing, walking across a room, or getting in or out of bed. Thus, state 2 is interchangeably referred to as the LTC or ADL state. The health state evolves according to a Markov process, where the probability matrix, $\pi_g(s'|t, s)$ is sex, age, and health-state dependent. Health status affects the distribution of out-of-pocket health expenditure shocks and the distribution of health status in the following year (including mortality risk). Out-of-pocket health expenditures (h) are lognormally distributed as a function of age, health status, and sex. Section V.A.2 describes the mapping of health variables between model and data.

Health-state-dependent and bequest utility.—In addition to affecting health costs and survival probabilities, health status affects preferences. There is a health-dependent utility function, such that spending when a consumer needs LTC is valued differently than spending when a consumer does not need LTC. When in good or poor health ($s \in \{0, 1\}$), consumers value consumption according to standard CRRA (constant relative risk aversion) preferences with parameter $\gamma > 0$:

$$U_{s \in \{0,1\}}(c) = \frac{c^{1-\gamma}}{1-\gamma}. \quad (1)$$

Utility when in need of help with LTC ($s = 2$) associated with chosen expenditure level c is

$$U_{s=2}(c) = (\theta_{\text{ADL}})^{-\gamma} \frac{(c + \kappa_{\text{ADL}})^{1-\gamma}}{1-\gamma}. \quad (2)$$

Upon death ($s = 3$), the agent receives no income and pays all mandatory health costs. Any remaining wealth is left as a bequest, b , which the consumer values with a warm-glow utility function:

$$v(b) = (\theta_{\text{beq}})^{-\gamma} \frac{(b + \kappa_{\text{beq}})^{1-\gamma}}{1-\gamma}. \quad (3)$$

When an individual is healthy or sick, utility is given by a power utility function of consumption. Bequests are valued using the standard warm-glow utility function developed by De Nardi (2004). When an individual needs LTC, utility is given by a similar formula, which treats LTC and bequests symmetrically in theory, allowing differences in preferences to be

determined empirically through estimated parameter differences.⁹ Two key parameters are θ and κ ; θ affects the marginal utility of an additional dollar spent, and κ controls the degree to which an expenditure is valued as a luxury good or a necessity, in the sense that it provides a utility floor or need. Since it is raised to the power $-\gamma$, increases in θ decrease the marginal utility of a unit of expenditure. The parameter κ allows the model to represent nonhomothetic preferences; an increase in κ indicates that the expenditure is valued as more of a luxury good; negative κ can be interpreted as the expenditure being a necessity.

C. Government

A person always has the option to use a means-tested government-provided care program. The cost of using government care is that a consumer's wealth is set to zero, while the benefit is that the government provides predetermined levels of expenditure, which depend on the health status of the individual. If a person chooses to use government care when not in need of LTC (i.e., when $s = 0, 1$), then the government provides a consumption floor; $c = \omega_G$, that is designed to represent welfare.

If an individual needs LTC ($s = 2$), then he or she must either purchase private LTC or use government care. Capturing the fact that LTC provision is essential for those in need and private LTC is expensive, there is a minimum level of expenditure needed to obtain private LTC, parameterized by χ ; that is, $c \geq \chi$ if $s = 2$ for those not using government care. In the model, government-provided care is loosely based on the institutions of Medicaid. If a person needs LTC and uses government care, the government provides $c = \psi_G$. The value ψ_G parameterizes the individual's value of public care, since that parameter determines the utility of an individual who needs LTC and chooses to use government care.

D. The Individual Problem

The individual takes r as given and chooses a' , c , and G to maximize utility. This problem, written recursively, is

⁹ We follow most papers in this literature, e.g., De Nardi, French, and Jones (2010, 2016a) and Lockwood (2018), by using the same exponent in the utility functions for all health states. We explore the sensitivity of our results to allowing for a bequest-specific exponent, γ_{beq} , in sec. IX. De Nardi, French, and Jones (2016a) also model a health-state-dependent utility function to study late-in-life saving patterns. Their focus is on medical spending more generally among a less wealthy population; they model an additively separable homothetic medical expenditure utility function with a marginal utility multiplier that varies as a function of nursing home need and other medical states (e.g., broken bones). By contrast, our health-dependent utility function is over total consumption in the LTC state and represents nonhomothetic preferences via κ_{ADL} . Thus, our parameterization provides the LTC utility function the same flexibility as the bequest utility function in terms of generating varied spending and saving behavior across wealth levels.

$$\begin{aligned}
V(a, y, t, s, h, g) = & \max_{a', c, G} \mathbb{I}_{s \neq 3} (1 - G) \{ U_s(c) + \beta E[V(a', y, t + 1, s', h')] \} \\
& + \mathbb{I}_{s \neq 3} G \{ U_s(\omega_G, \psi_G) + \beta E[V(0, y, t + 1, s', h')] \} \\
& + \mathbb{I}_{s=3} \{ v(b) \}, \text{ subject to} \\
a' = & (1 - G)[(1 + r)a + y(t) - c - h] \geq 0, \\
c \geq & \chi \text{ if } (G = 0 \wedge s = 2), \\
c = & \psi_G \text{ if } (G = 1 \wedge s = 2), \\
c = & \omega_G \text{ if } (G = 1 \wedge (s = 0 \vee s = 1)), \\
b = & \max\{(1 + r)a - h', 0\}, \\
U_s(c) = & \mathbb{I}_{s \in \{0,1\}} \frac{c^{1-\gamma}}{1-\gamma} + \mathbb{I}_{s=2} (\theta_{\text{ADL}})^{-\gamma} \frac{(c + \kappa_{\text{ADL}})^{1-\gamma}}{1-\gamma}, \\
v(b) = & (\theta_{\text{beq}})^{-\gamma} \frac{(b + \kappa_{\text{beq}})^{1-\gamma}}{1-\gamma}.
\end{aligned}$$

The value function has three components, corresponding to the utility plus the expected continuation value of a living individual who chooses not to use government care, that of one who chooses to use government care, and the warm-glow bequest utility of the newly deceased individual.¹⁰ The parameter $G = 1$ if the consumer chooses to use government care, and $G = 0$ if the consumer chooses not to use government care. A person using government care has expenditure levels set to predetermined public care levels and zero next-period wealth. The budget constraint shows that wealth next period is equal to zero if government care is used, and it is otherwise equal to the return on savings plus income minus chosen expenditures minus the health cost shock. The individual cannot borrow and cannot leave a negative bequest, and private expenditure when in need of LTC must be at least χ .

E. Describing Optimal Behavior

In this section, we explore key properties of optimal individual behavior to illustrate how each force in the model contributes to consumption and savings patterns over the life cycle and across the income and wealth distributions. The individual's saving behavior is largely determined by the confounding influence of the precautionary saving motive and the bequest motive in the presence of government policies. LTC needs occur

¹⁰ Technically, there is a fifth health state that is reached (with certainty) only in the period after death and is the absorbing state, so that the consumer receives only the value of a bequest in the first period of death.

with nontrivial probability, and paying for such care privately is very costly. The fact that the government offers a means-tested public care option induces interesting behavior. Because the individual has the option to choose government care, the value function is nonconcave, and the optimal saving policy is discontinuous. The model does not permit analytic solutions and must be solved numerically, with details of our solution algorithm presented in Ameriks et al. (2018a).¹¹

1. Discontinuous Saving Policies

The option to use means-tested government care induces discontinuous saving policies, in a manner similar to that studied by Hubbard, Skinner, and Zeldes (1995). Roughly speaking, very high-wealth individuals have enough savings to ensure they will obtain a high level of personal consumption and leave a large bequest, regardless of whether or not they need to pay for private LTC. Low-wealth individuals, even if they saved almost all of their money and consumed very small amounts each year, would not be able to save enough to make it optimal for them to purchase private LTC if they eventually needed it. Thus, it is the middle-wealth people—like those in the VRI—whose actions are most likely to be affected by precautionary saving motives. If these middle-wealth individuals are frugal and save, they will have enough wealth to purchase private LTC if they need it late in life. If they do not save, but rather consume at a high rate over their life cycle, they will have higher utility when alive and healthy but will forgo a bequest and rely on public provision of LTC if they need it later in life. There exists some threshold wealth level, conditional on all other idiosyncratic state variables, such that it is optimal for all agents with more wealth to follow the frugal path and for all agents below to follow the spendthrift path, with a discrete difference in their saving policy for a tiny difference in their wealth state. To illustrate optimal consumer behavior, we present model simulations at certain parameter values.¹² Parameters are estimated and discussed further in section VI.B.

The discontinuity of the saving policy is demonstrated in figure 1 by plotting the objective function that corresponds to the nonoptimized

¹¹ The nonconcavity of the value function and the discontinuity in the optimal savings policy introduce computational complications. We use a modified endogenous-grid method, building on insights from Fella (2014). The model solves approximately 10 times faster when using the modified endogenous-grid algorithm, compared to value function iteration, which is essential, since estimation of the model requires computational efficiency.

¹² Parameter and state variable values for fig. 1 were chosen to clearly illustrate the discontinuity in the saving policy, which is relevant most often for households with low levels of assets. All other figures in this section are plotted at baseline parameter values or explicitly compare baseline outcomes to those generated by alternative parameter values.




FIG. 1.—Saving policy discontinuity. a = assets in thousands of dollars.

value function across saving policies for different initial levels of wealth. That is, plotted on the horizontal axis is age $t + 1$ wealth, and on the vertical axis is the value associated with that amount of saving. The three lines depict the graph for an individual with identical states aside from the three different age- t wealth levels: \$33K (dotted line), \$36K (solid line), and \$39K (dashed line). In this example, the optimal value of saving for the low-wealth individual is zero. This person is forgoing all precautionary saving and consuming as much as possible today. The higher-wealth individual has enough money that it is worth it to save to help smooth consumption across states and time. This can be seen in the top line, in which the maximum of the value function is achieved with savings around \$32K, attaining a noticeably higher value than if this individual saved zero. As presented in the middle line, there exists an age- t wealth level for which the global maximum value jumps from the lower to the higher-savings local maximum. An individual with \$36K in this example is nearly indifferent between saving around \$31K and saving zero. It is around this level of current wealth that there is a discrete jump in the optimal saving policy (the value function is kinked but remains continuous).

2. How Preferences and States Determine Saving Behavior

Saving decisions are ultimately determined by the preferences of individuals and by their environment. As was highlighted by Dynan, Skinner, and Zeldes (2002), a dollar saved today is fungible in its future use. Saving early in the person's life could be to insure against future uncertain events such as LTC as well as to ensure that suitable savings remain at end of life to leave a desired bequest. If the bequest motive is weak, oversaving for an uncertain late-in-life event that never occurs is costly, as the individual would much rather have had a smooth higher-consumption path over his life. However, with a strong bequest motive, "extra" savings at the end of life are highly valued, which reduces the cost of ex post oversaving.

To demonstrate how savings are influenced by bequest and LTC-induced saving motives, we plot various age profiles for wealth and expenditure for a simulated individual, in response to different sequences of health shocks, for different initial states, and for different preference parameters. Unless otherwise stated, the figures plot the wealth and expenditure profiles of a male who starts healthy at age 55 and has the median income profile, median wealth, and preference parameters from our preferred baseline estimation. Unshaded areas indicate behavior when healthy ($s = 0$), while shaded regions indicate behavior when in need of LTC ($s = 2$). It is important to note that these patterns will not be representative of wealth and expenditure profiles of the population, as these are individuals and shocks selected to illustrate the workings of the model and are not necessarily typical or representative of the VRI sample or the US population.

As a baseline, figure 2 shows the wealth and consumption paths of a man who receives a shock sequence such that he remains in good health until death at $T = 108$. Wealth accumulates until age 75 and then steadily decumulates with age. Early on, the individual saves, driven by a combination of LTC and bequest motives. As the individual ages, the probability of needing LTC for any given year tends to increase, but eventually the chance that LTC will be needed for any given large number of years decreases. The increase in consumption with age occurs because the individual was saving precautionarily for LTC, and as he continues to receive such a good run of positive health shocks, he starts to consume the ex post extra savings slowly.

Figure 3A demonstrates the rapid dissaving and high expenditure associated with the need for LTC. This person received health shocks such that he was healthy his entire simulated life, except for one period in which he needed LTC for ages 74–76, highlighted by the shaded region. At the onset of needing LTC, expenditures jump from around \$60,000




FIG. 2.—Wealth and expenditure profiles for healthy males.

per year to around \$120,000 per year, resulting in a large decrease in wealth. Expenditure remains high and roughly constant during the 3-year LTC period, as savings decline rapidly. After 3 years of LTC, the individual steadily dissaves and consumes, as no other adverse health shocks occur until death. Saving and expenditure behavior depend on an individual's level of wealth. Figure 3*B* plots the behavior of an individual who is similar in all ways except for having lower wealth at age 55. The low-wealth individual saves more aggressively early on in order to build a buffer stock of wealth in case LTC is needed and in order to be able to leave a bequest. Similar patterns of rapid dissaving and high levels of expenditure are associated with the LTC event. However, the low-wealth individual actually increases wealth after exiting LTC to return to a desired buffer-stock level of wealth.

Figure 4*A* documents the behavior of a lower-wealth individual who also has the lower first-quintile income profile. Furthermore, compared to figure 3, in this simulation his need for LTC starts earlier and lasts longer. At first he purchases private LTC, but the high level of expenditure associated with his need for LTC depletes his wealth to near zero, at which point he chooses to use publicly provided LTC for the rest of his LTC episode and then live hand to mouth afterward. Note that public care expenditure (dashed line) is included in the total expenditure reported.¹³ Figure 4*B* shows what happens if the individual started at age 55 with \$30,000 in savings instead of \$100,000. He consumes very

¹³ For a discussion of the level of public care expenditure (ψ_G), see sec. VI.B.

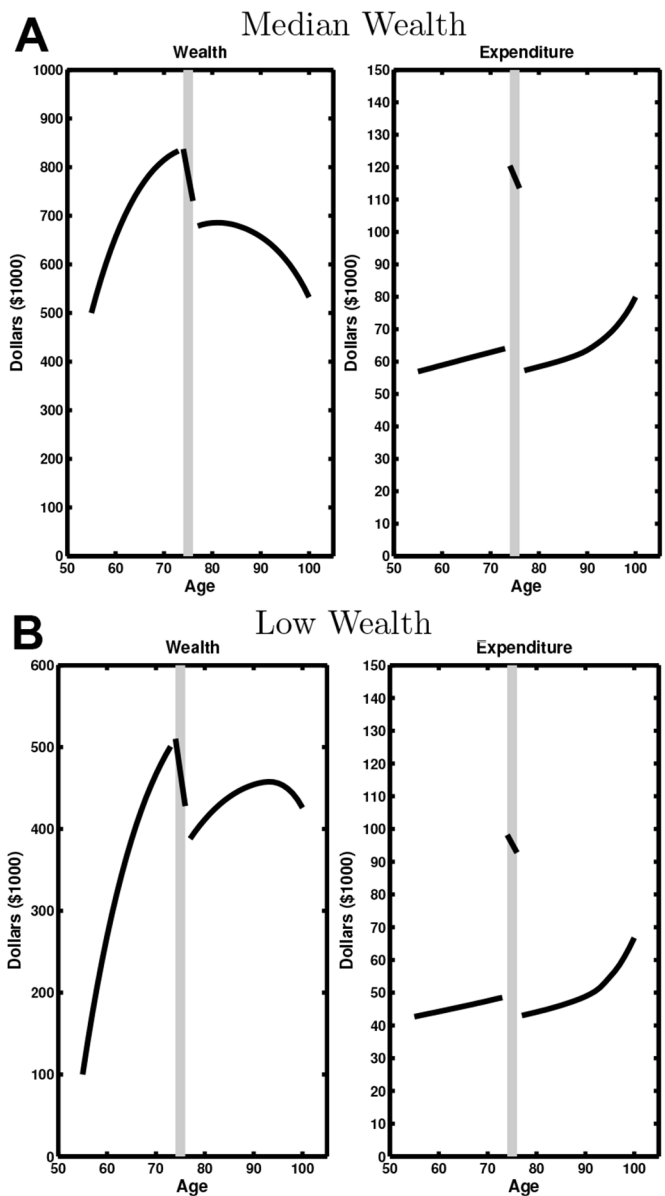


FIG. 3.—Wealth and expenditure profiles for median-income males.

little and saves up until he needs LTC. His wealth is so low that he immediately uses public care as soon as he needs LTC. When he no longer needs public care, he simply consumes his roughly \$20,000-per-year income. As is apparent, the need for extended LTC rapidly depletes

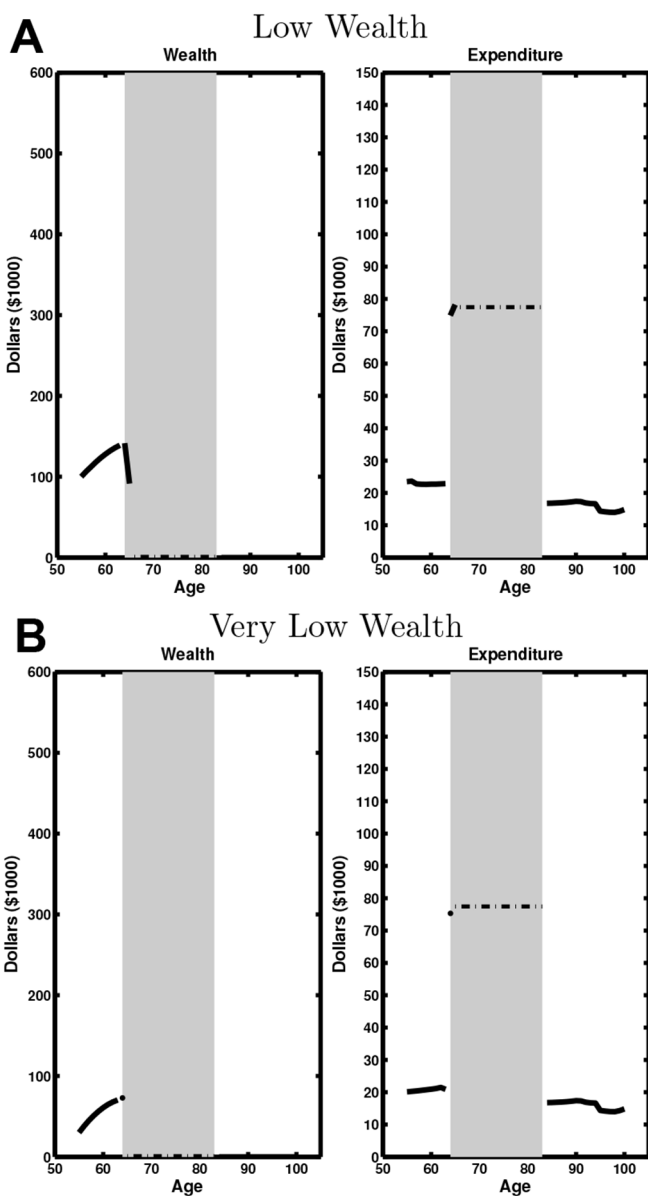


FIG. 4.—Wealth and expenditure profiles with publicly provided LTC.

savings and can lead to extended periods of reliance on government care and low consumption for the remainder of life.

Quantitatively, the levels of expenditure are quite reasonable across the wealth and income distribution. An individual who has \$800,000 in wealth at age 74 and earns around \$50,000 a year spends around \$120,000 if he needs LTC that year, while an individual who has \$150,000 in wealth at age 74 and earns \$20,000 a year spends around \$75,000.

These saving and expenditure patterns are strongly influenced by people's preferences. To demonstrate the importance of the health-state utility function, figure 5A recreates the simulation presented in figure 3A, except for an individual with preferences such that spending when in need of LTC is valued just as spending when healthy ($\theta_{ADL} = 1$, $\kappa_{ADL} = 0$). The original behavior induced by baseline preferences is drawn with dotted lines, and that associated with alternative preferences is drawn with solid lines. This analysis shows that much of the increase in wealth during the individual's 50s and 60s was driven by the precautionary saving motive associated with LTC. Furthermore, expenditure levels when in need of LTC are much closer to expenditure when healthy, with a slight uptick due to the increased mortality risk associated with the worse health state. This major change in expenditure patterns foreshadows that our estimated health-state utility function induces higher marginal utility of expenditure when in need of LTC, not less.

Both the health-state utility function and the bequest function affect saving and spending behavior. Figure 5B plots the life-cycle behavior of the same median-wealth, median-income individual, but with parameters such that bequests are more strongly valued. As can be seen, the stronger bequest motive increases savings early on. Furthermore, the stronger bequest motive has a significant effect on late-in-life wealth levels, leading an individual to reach age 100 with over 70% more wealth than the baseline individual. This person needed to save so much early on because he had a strong desire to spend when in need of LTC and to leave a bequest. Later in life, expenditure patterns look similar, because consumption similar to that in the baseline case can be sustained without depleting wealth, as a result of the higher level of financial income generated by a larger stock of wealth.

Finally, to give a sense of the effect of health-state utility on saving and bequest behavior in the data, we plot in figure 6A a more typical health pattern, in which a median-wealth, median-income man is healthy from age 55 to 84, needs LTC for age 85–89, and then dies. The figure plots wealth over the life cycle with the baseline parameters and also with no-ADL-state utility. Without ADL-state utility, this person accumulates wealth roughly from age 55 to 65 and then starts steadily dissaving. When in need of LTC, his rate of dissaving is only marginally faster, driven by changes in life expectancy. In contrast, with the baseline ADL-state utility

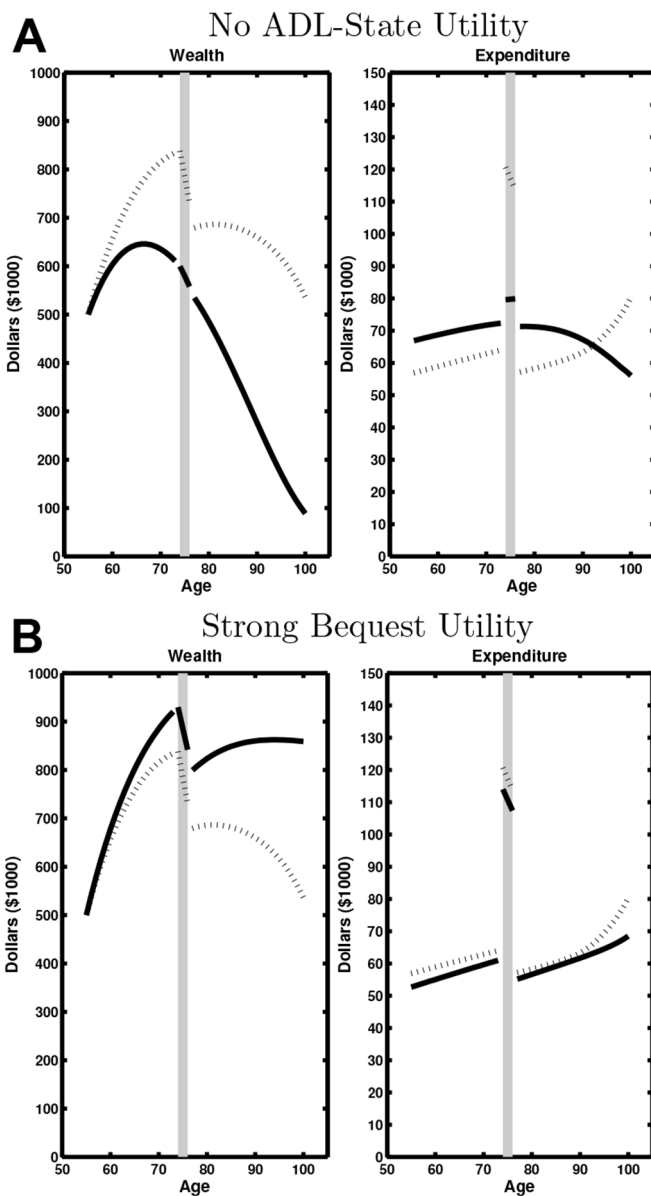


FIG. 5.—Wealth and expenditure profiles with baseline (dotted line) and alternative (solid line) preferences.

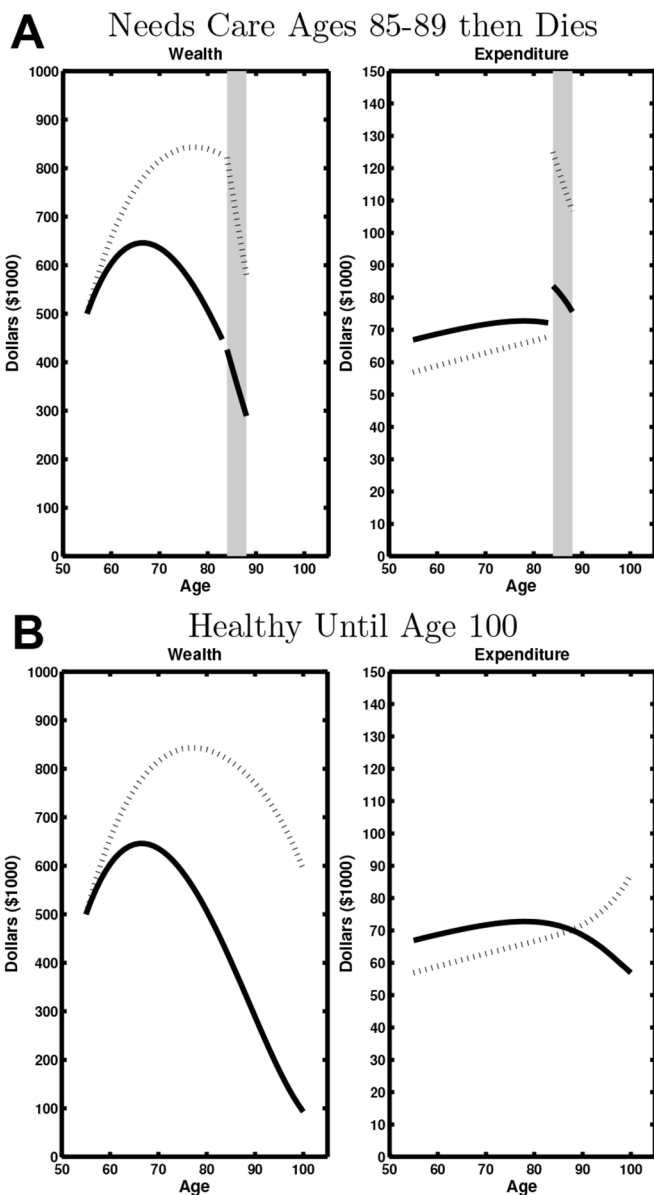


FIG. 6.—Wealth and expenditure profiles with (dotted line) and without (solid line) health-state utility for different health paths.

function, this person accumulates wealth more rapidly from age 55 to 65 and continues to accumulate wealth until age 80. Furthermore, when the negative health shock realizes, dissaving is rapid, driven by the higher utility of expenditure in this health state. This desire to spend when in need of help with ADLs is, to a large degree, why he was saving in the first place. In this scenario, the baseline parameterization results in almost double the bequest, compared to the no-ADL-state utility case, even with the rapid health-related spend-down at the end of life due to the increased precautionary saving earlier in life. Many people die without ever needing LTC. Figure 6*B* plots savings over the life cycle for a man who is healthy until death at age 100, with and without the health-state utility function. This figure demonstrates that the health-related precautionary saving motive generates a large incidental bequest even though there is the typical spend-down late in life due to mortality risk, absent LTC needs.

With an understanding of the key features of optimal saving behavior in the model and how they relate to important state and parameter values, we turn to a description of the data, with which these parameter values will be estimated.

III. Financial and Demographic Data

In order to examine late-in-life wealth patterns, it is essential to have data on a population with large enough financial resources to face nontrivial spending, saving, and giving decisions. This paper draws on the newly developed VRI, which combines survey and administrative account data. In this section, we briefly describe the VRI, highlighting the advantages of the sample population for addressing the question at hand.

The VRI consists of approximately 9,000 individuals drawn from Vanguard account holders who are at least 55 years old. Additionally, we require Vanguard assets of at least \$10,000 (to assure nontrivial engagement with Vanguard) and internet registration with Vanguard (to allow for surveys administered over the internet). As a point of comparison, the VRI is cross sectionally about the same size as the HRS and around four times larger than the Survey of Consumer Finances (SCF) in the relevant age group. Surveys are administered over the internet and ask respondents about their and their spouse's or partner's wealth, income, and decision-making motives.

A sample drawn from Vanguard account holders is, of course, not random or representative of the US population. For example, by design, the sample is drawn from individuals who have positive financial wealth. Hence, we exclude the large fraction of households who approach or reach retirement age with few or no financial assets. The population with little financial wealth, by construction, does not face a wealth-allocation problem for LTC. Use of this new data set is a significant contribution of

this paper. It provides a large sample of older Americans with sufficient financial assets to face meaningful trade-offs between consumption across time, between spending when well and spending when in need of assistance, between LTC in private and that in publicly funded facilities, and between leaving bequests and spending while alive. We also replicate the main results, using data from the HRS, which provides a representative sample in an age range similar to that in the VRI. Using estimates based on responses from the VRI, we are able to confirm that the predictions of the model are broadly consistent with the wealth moments from the HRS. This finding suggests that these results can be generalized beyond populations sharing the demographic and economic characteristics of the VRI.

Since we do not explicitly model the family, in this paper we restrict our data to include only single respondents, who were oversampled to ensure a large single subsample. For the remainder of this paper, we focus on a sample of 1,241 singles with no missing survey responses to mandatory questions.

For points of comparison, we construct “VRI-eligible” subsets of the HRS and the SCF by imposing sample screens to parallel the VRI: age 55 years and older, financial assets of at least \$10,000, and access to the internet. After imposition of these screens, the characteristics of the VRI sample are similar in many dimensions to these subsets of the 2012 HRS and 2013 SCF, representing individuals in roughly the upper half of the wealth distribution. Financial wealth is defined as the sum of an IRA (individual retirement account) and employer-sponsored retirement, checking, saving, money market, mutual fund, certificate of deposit, brokerage, and education-related accounts, plus the current cash value (if any) of life insurance and annuities. Income is defined as the sum of labor income, publicly and privately provided pensions, and disability income. Table 1 compares wealth and income of the VRI and VRI-eligible subsets of the HRS and SCF restricted to the single households considered in this paper. Our sample is well positioned to complement existing samples with a highly relevant population. In table 1, we see that VRI respondents have more wealth than the VRI-eligible HRS and SCF respondents but that the differences are much less stark than those with the overall population, which has close to zero wealth at the median. Furthermore, although the income is somewhat higher in the VRI than in the VRI-eligible HRS, the VRI and the VRI-eligible SCF have very similar levels of income. A downside of using VRI data is that we have only a cross section of measurements, which does not allow us to adjust for cohort or time effects and may reflect mortality bias.

For more details, we refer the reader to Ameriks et al. (2015), which provides an exhaustive analysis of the VRI, both on the survey methodology, which measures wealth from an accounts-based construct, and

TABLE 1
INCOME AND WEALTH DISTRIBUTION ACROSS SURVEYS: VRI-ELIGIBLE SINGLE HOUSEHOLDS

SURVEY	OBSERVATIONS	MEAN	PERCENTILE				
			10th	25th	50th	75th	90th
Financial Wealth							
VRI	1,241	808,007	101,000	262,113	527,600	993,800	1,602,000
HRS	1,201	376,432	24,000	68,000	178,000	445,000	920,000
SCF	265	487,234	18,500	58,500	159,000	410,700	1,019,000
Income							
VRI	1,241	69,452	17,500	34,223	56,000	86,550	121,473
HRS	1,201	65,402	10,860	18,817	36,000	65,000	105,012
SCF	265	80,963	25,363	35,509	51,741	85,221	121,744

NOTE.—Financial wealth is the sum of an IRA and employer-sponsored retirement, checking, saving, money market, mutual fund, certificate of deposit, brokerage, and education-related accounts plus the current cash value (if any) of life insurance and annuities. Income is defined as the sum of labor income, publicly and privately provided pensions, and disability income. The sample comprises single households meeting VRI sample screens: age 55 years and older, assets of at least \$10,000, and internet access. Income is total household income (excluding distributions from defined-contribution pension plans). Data are from the 2014 HRS and the 2013 SCF. See apps. B and C of Ameriks et al. (2015) for a discussion of the definitions of variables in the HRS and SCF and for a detailed comparison of the VRI, HRS, and SCF.

on the resulting collected data. For the purposes of this paper, it is most important to note that the VRI contains high-quality measures of individuals’ wealth and income and, crucially, responses to SSQs that were specifically designed to identify parameters of the model just developed. The VRI also has measures of self-reported health status and need for help with ADLs, elicited using the same questions as in the HRS, which we use to estimate health-state transition matrices. To estimate both the health-state transition probabilities and the out-of-pocket health expenditure shock distributions, we use HRS panel data, as detailed in section V.

IV. SSQs

SSQs are stated-preference questions, with scenarios that are quite closely linked to important life decisions faced by our respondents. Behavior in the model is driven by the preferences of individuals and the economic environment in which they make choices. Since a main goal of this paper is to identify the relative contributions of different saving motives associated with different preferences, it would be ideal if survey respondents could accurately and directly report their preference parameters. Of course, we cannot ask survey respondents to report their coefficient of relative risk aversion, much less θ_{ADL} . Thus, if we want to develop direct measures of preferences, we need to develop survey instruments that

allow respondents to provide us with information that identifies preference parameters not only in a language in which they are comfortable but also in a format that allows a precise mapping to structural parameters of interest.

Along these lines, revealed-preference methodology uses observed choices to perform inference about preferences. If a utility function is assumed to represent preferences, often these observed behaviors can be used to estimate preference parameters. In a similar vein, we develop SSQs that use choices made in hypothetical scenarios to estimate preference parameters. In constructing our survey, we create a highly structured hypothetical environment with a very restricted choice set that allows us to make fewer assumptions on the unspecified economic environment. Though necessarily incomplete *per se*, our scenarios are significantly more detailed than those in typical hypothetical questions. Our questions are designed to provide the survey respondent precise details on all relevant individual states of the world from the perspective of the structural model. SSQs ask the respondent to comprehend and imagine complex scenarios. As with any hypothetical question, there are legitimate concerns about whether survey respondents can understand the scenario and whether they can respond from the perspective of that scenario. To make these tasks as easy as possible for the survey participant, we paid close attention to the presentation of the material and developed the survey with input from survey design experts and cognitive psychologists in a process of exploratory cognitive interviews, pilot surveys, and then production surveys; we also performed tests of respondent comprehension and of the coherence and consistency of responses to address such concerns. See section IV.C for analysis of SSQ responses and Ameriks et al. (2018c) for further evidence on the credibility and coherence of SSQ responses.

The SSQs are one ingredient in our overall estimation strategy to identify preference parameters. We treat each individual as an optimizer in a specific problem and characterize the individual by financial and demographic variables and preference parameters. The SSQs ask people to make choices that would be very revealing about their preference parameters if only we were to observe them making such choices. A key feature of the SSQ approach is that it provides data about the priority of LTC risks and desires for all individuals in the sample, not just the ones who end up needing LTC. We add as moments in the GMM (generalized method of moments) estimation the mean of each SSQ, so that we are estimating preferences of an optimizer who would answer the SSQs in a manner consistent with the central tendencies of how the survey respondents answered the SSQs. Because these are deep structural preference parameters, they also affect all other behavior of these modeled individuals, including their saving behavior. In addition to the SSQs, we

include wealth moments conditional on age to make sure that our optimizing individual has preferences that are also consistent with the saving behavior of our survey respondents. There is a weighting matrix that determines the weight on the different moments, but all moments are informative of all preference parameters. In the context of this study, it turns out that two very different types of individuals could equally well have generated the wealth data: one with strong bequest motives and weaker LTC-related spending desires and vice versa. The SSQs allow us to sharply identify behavior that is only weakly identified in the wealth data alone: an individual with a very strong active bequest motive and weak LTC utility would not answer the SSQs as the survey respondents did, while an individual with strong LTC utility and a somewhat weaker active bequest motive would answer the SSQs like our survey respondents.

A. *Detailing SSQ 2: Spending in Good Health versus When Needing LTC*

Ultimately, the model is estimated on responses from four types of SSQs and wealth data. In this section, we first illustrate the key features of SSQs by detailing one particular SSQ related to LTC (SSQ 2), and then we present the other three SSQs. In section V, we detail more precisely how we construct moments and use model simulations to estimate parameters.

In SSQ 2, we are interested in understanding how individuals trade off having wealth in states of the world when they do not need LTC and when they do need LTC. At the core of the question, we are asking individuals to solve a simple portfolio-allocation problem. The researchers specify that the respondent has some wealth (W), faces some chance that they will need LTC ($1 - \pi$) and some chance that they will not need LTC (π), and must allocate their resources by purchasing state-contingent assets (z_1, z_2), given a relative price of $z_2(p_2)$, to finance expenditure in the two possible states of the world. In the survey, we set $p_2 = 1/(1 - \pi)$. The mathematical representation of the survey question that we use for estimation is

$$\begin{aligned} \max_{z_1, z_2} \quad & \pi \frac{z_1^{1-\gamma}}{1-\gamma} + (1-\pi) \frac{(\theta_{\text{ADL}})^{-\gamma} (z_2 + \kappa_{\text{ADL}})^{1-\gamma}}{1-\gamma}, \\ \text{subject to} \quad & z_1 + p_2 z_2 \leq W, \\ & z_1, z_2 \geq 0, \text{ and } z_2 \geq -\kappa_{\text{ADL}}. \end{aligned} \tag{4}$$

1. Identification

The first-order condition of the optimization problem gives the optimal allocation as a function of preference parameters. By inverting this

function, we map the allocations chosen by survey respondents to preference parameters. For example, the optimal decision rule of the above problem is given by

$$z_2 = \begin{cases} 0 & \text{if } \pi \left(\frac{W}{p_1} \right)^{-\gamma} \left(\frac{p_2}{p_1} \right) - (1 - \pi)(\theta_{\text{ADL}} \kappa_{\text{ADL}})^{-\gamma} > 0, \\ \frac{W[\pi p_2/p_1(1 - \pi)]^{-1/\gamma} - p_1 \theta_{\text{ADL}} \kappa_{\text{ADL}}}{p_1 \theta_{\text{ADL}} + p_2[\pi p_2/p_1(1 - \pi)]^{-1/\gamma}} & \text{otherwise.} \end{cases}$$

The scenario exogenously specifies π and W . By obtaining responses for two different combinations of W and π for which allocations to both states are positive, we are able to identify θ_{ADL} and κ_{ADL} , given γ . Combining these with the results of SSQ 1, which identify γ , we are thus able to identify all three parameters from responses to these questions.

The key survey design challenge is that most individuals cannot understand the allocation problem in the mathematical language of optimal control. We present below the SSQ that is designed to help survey respondents provide (z_1, z_2) such that they are making a choice that we are confident corresponds to that in the optimization problem, but in a format in which they are capable of doing so. Identification for all preference parameters proceeds similarly, as demonstrated below.

2. Survey Design

To ease respondent comprehension, these questions are presented in five steps. In the first screen, we begin by telling the respondent explicitly what trade-off we are asking them to think about. This is done to prompt the respondent to weigh the relevant risks we are interested in and to alleviate their concern over not understanding the point of the question and guessing about the motives of the survey designers. Second, the question presents the specific scenario and details the choices that the respondent must make. This screen is the complete scenario and is made available to the respondents as they are giving their final answers if they would like to check any features of the scenario. Third, we present a set of rules that further defines the environment and clarifies the span of options and resources available to the survey respondent in each scenario. Fourth, we verify comprehension and reinforce key features of the question with a set of multiple-choice questions about the scenario. Finally, we record an answer, using a custom-designed interactive slider. See Ameriks et al. (2018c) for a detailed discussion of the SSQ

design process and further analysis of individual responses that demonstrate the credibility of SSQs.

3. The Scenario

The survey instrument first states the scenario precisely but as simply as possible, consistent with being precise. Specifically, the survey displays a screen with the following text.¹⁴

We are interested in how you trade off your desire for resources when you do and when you do not need help with activities of daily living (ADLs). This scenario is hypothetical and does not reflect a choice you are likely ever to face.

Suppose you are 80 years old, live alone, rent your home, and pay all your own bills. Suppose that there is a chance that you will need help with ADLs in the next year. If you need help with ADLs you will need long-term care.

- There is a **25%** chance that you **will** need help with ADLs for all of next year.
- There is a **75%** chance that you **will not** need any help at all with ADLs for all of next year.

You have **\$100,000** to divide between two plans for the next year. This choice will affect your finances for next year alone. At the end of next year you will be offered the same choice with another \$100,000 for the following year.

- Plan C is hypothetical ADL insurance that gives you money if you **do** need help with ADLs.
 - For every \$1 you put in Plan C, you will get **\$4** to spend if you need help with ADLs.
 - From that money, you will need to pay all your expenses including long-term care at home or in a nursing home and any other wants, needs, and discretionary purchases.
- Plan D gives you money only if you **do not** need help with ADLs.
 - For every \$1 you put in Plan D, you will get **\$1** to spend if you **do not** need help with ADLs.
 - From that money, you will need to pay for all of your wants, needs, and discretionary purchases.

¹⁴ In previous sections of the survey, the definition of “needing help with ADLs” is given, and understanding is verified. Further, a reminder of this definition appears if respondents move their mouse over the word “ADLs” in the scenario.

4. Presenting the Rules of the Scenario

Immediately after the scenario is presented, the respondents are provided with a recap and elaboration of the specific rules that govern their choice. This recaps the previous screen but is presented in a bulleted, easy-to-read format. In addition, some features that were hinted at in the first screen—for example, that there is no public care option and that determination of which plan pays out is made by an impartial third party—are stated explicitly. These rules are designed to ensure that the word problem corresponds as closely as possible to the intended optimal control problem.

- You can only spend money from Plan C or Plan D next year. You do not have any other money.
- If you want to be able to spend whether or not you need help with ADLs, you need to put money into both plans.
- If you need help with ADLs, all money in Plan D is lost.
- If you do not need help with ADLs, all money in Plan C is lost.
- Any money that is not spent at the end of next year cannot be saved for the future, be given away, or be left as a bequest.
- You must make your choice before you know whether you need help with ADLs. Once you make your choice, you cannot change how you split your money.
- Regardless of whether or not you need help with ADLs, your hospital, doctor bills, and medications are completely paid by insurance.
- Other than Plan C, you have no other resources available to help with your long-term care. **You** have to pay for any long-term care you may need from Plan C.
- There is **no public-care option or Medicaid** if you do not have enough money to pay for a nursing home or other long-term care.
- An impartial third party that you trust will verify whether or not you need help with ADLs immediately, impartially, and with complete accuracy.

5. Verification Questions

In order to reinforce details of the scenario and measure comprehension, we ask the respondents a sequence of questions about the specifics of the scenario, including payoffs in different states, potential uses of money, potential expenses, and rules regarding the payouts. When answering these questions, the respondents do not have access to the screens describing the scenario, but they have a chance to review the information before retrying any missed questions a second time. If the respondents fail

to answer questions correctly a second time, they are presented the correct answers. As documented in Ameriks et al. (2018c), the vast majority of individuals answered almost all verification questions correctly before recording their responses to the SSQs.

6. Recording the Response

Having reinforced and measured understanding, we are finally prepared to ask the question: how would respondents split their wealth between the two plans? After again presenting them with the original scenario, we present them a screen—a snapshot is shown in figure 7—with a link in the top-right corner to the full scenario. The responses are recorded through an interactive slider that we developed for the purpose of eliciting responses to SSQs. The slider allows the respondent to experiment with different answers and dynamically displays the trade-offs implicit in the SSQs—in this case the trade-off between spending when well and spending when needing LTC. The axis is not labeled with dollar amounts. Instead, the screen contains indications that moving the slider right places more money in plan D and moving it to the left places more money in plan C. These amounts are displayed dynamically at the ends of the slider. See http://ebp-projects.isr.umich.edu/VRI/survey_2.html for the SSQ survey instrument, including the slider.

[Click here for complete scenario](#)

Please make your decision on splitting money into Plan C and Plan D by clicking on the scale below. To put more money in Plan C, move the slider to the left. To put more money in Plan D, move the slider to the right. The numbers in the boxes will change as you move the slider to let you know how much you will receive if you need long term care and if you do not.

Please move the slider to see how it works. When you are ready, place the slider at the split you want and click NEXT to enter your choice.



FIG. 7.—SSQ 2 response screen.

This mode of presentation likely contributes to the high-quality responses we were able to elicit to the SSQs. Use of the slider is effective because it summarizes the key trade-off in each of the SSQs and allows the respondent to see how the trade-off operates. The implementation of the slider also addresses several issues with surveys in general. First, there is no value on the slider when it is first presented. The respondent must click to establish an initial value. Second, the survey asks the respondent to move the slider from this initial point (and indeed requires that they do so before recording a response). Together, these two features of the presentation serve to reduce the effect of anchoring. Indeed, we observe little effect of a respondent's first click on their final answer.

After recording a response to this initial question, we ask two variations of this SSQ with different wealth levels and probabilities of needing LTC. This provides further information about how they value having wealth in different states and provides us with a consistency check of individual choices.

B. Overview of the Other SSQs

In addition to the SSQ presented above, which examines the trade-off between wealth when in need of assistance with ADLs and that when not in need, the survey presents three other SSQs that examine trade-offs that are relevant to understanding late-in-life saving motives. Like the above-presented SSQ, these questions present situations that individuals may be unlikely ever to face. Yet they capture the same type and scale of risks and trade-offs that individuals do need to confront in making lifetime saving decisions. The choices that individuals make when confronted with these hypothetical situations provide information regarding the relative values of having wealth in different states of the world. These three additional SSQs are outlined below, and the full text of all four SSQs, as well as histograms of the survey responses, are available in Ameriks et al. (2018b).

1. SSQ 1: Risk Aversion

The first type of SSQ posed is a modified version of the Barsky et al. (1997) question (BJKS question) that examines an individual's willingness to trade a certain lifetime income for a lottery over lifetime income that has a higher expected payoff. Their original question measured tolerance for risk and has been used frequently to identify the coefficient of relative risk aversion parameter in a power utility function. In the VRI formulation, we refine this question by specifying a more precise environment in which age, health expense, labor income, unexpected expenses,

and outside sources of wealth are all controlled for. More generally, the framing of SSQ 1, relative to that of the BJKS question, represents an evolution of thought favoring being more explicit about the precise environment of SSQ scenarios at the expense of having longer and more complex survey questions. We also make the decision a (repeated) static choice, by allowing the individuals to bet only over a single year's spending at one time. This significant departure from the standard Barsky et al. formulation is necessary to avoid confusion with late-in-life health-state utility and bequest preferences.

Specifically, in the VRI question, we present individuals an option of choosing between two plans that affect their consumption for the upcoming year. The first plan guarantees \$100,000 for certain, and the second plan will, with 50% probability, double income to \$200,000 or, with 50% probability, reduce income by some fraction. The individuals are asked a series of questions that categorize them into ranges of proportional losses that they would be willing to risk, before finally being prompted to provide a point estimate of the largest possible loss for which they would choose the risky lottery over the certain income option. SSQ 1 follows the BJKS question by asking about preference over discrete gambles. At the end of the sequence, SSQ 1 uses plain language to ask respondents to provide the ξ that satisfies

$$\frac{W^{1-\gamma}}{1-\gamma} = 0.5 \frac{(2W)^{1-\gamma}}{1-\gamma} + 0.5 \frac{[(1-\xi)W]^{1-\gamma}}{1-\gamma}, \quad (5)$$

with the choice of ξ point-identifying γ . A variant of this question is then presented, comparing lotteries to \$50,000 in certain income.

2. SSQ 3: Bequest versus LTC Utility

In the third type of SSQ that we posed (the second being described in sec. IV.A), we ask individuals to make an irreversible portfolio decision that allocates money between bequests and expenditure while alive when the individuals need help with ADLs. This question, which is similar to one posed in Ameriks et al. (2011), removes the possibility of an incidental bequest and thus allows us to focus on an intentional bequest motive. Because bequests observed in standard data sources also include unused precautionary savings, it is difficult to identify how strong the bequest motive is. By removing the option of saving money usable for both precautionary and bequest purposes, we are able to separately identify the relative strength of the two motives. To formulate this question, we present individuals with \$100,000 and tell them that they have exactly one year left to live. Furthermore, they will need help with ADLs for the entire year. They then must allocate money between two plans: one that is available

for them to spend during the coming year but cannot be left as a bequest and another that is accessible only as a bequest upon their death. This response is then recorded, and the individuals are asked how their portfolio allocation would change if they had \$150,000 or \$200,000 of wealth. SSQ 3 maps to the following optimization problem:

$$\begin{aligned} \max_{z_1, z_2} & \frac{(\theta_{\text{ADL}})^{-\gamma}(z_1 + \kappa_{\text{ADL}})^{1-\gamma}}{1 - \gamma} + \frac{(\theta_{\text{beq}})^{-\gamma}(z_2 + \kappa_{\text{beq}})^{1-\gamma}}{1 - \gamma}, \\ \text{subject to } & z_1 + z_2 \leq W, \\ & z_1, z_2 \geq 0, z_1 \geq -\kappa_{\text{ADL}}, \text{ and } z_2 \geq -\kappa_{\text{beq}}. \end{aligned} \quad (6)$$

3. SSQ 4: The Value of Public Care

In the final SSQ, we focus on an individual's willingness to utilize public LTC. The environment is similar to that of SSQ 3, in that the respondents are told they have only one year to live, that will need help with ADLs for the entire year, and that the only two spending channels accessible to them are spending on themselves during the year and leaving money as a bequest. In this scenario, there is a publicly funded care option that is available to them. Using the public care option will allow them to leave all of their wealth as a bequest, but they will receive the level of care that a typical public care facility would provide. We then ask for the level of wealth at which they would be indifferent between taking public care and paying for their own. Intuitively, for extremely low levels of wealth, the respondents are likely to utilize public care, as they are unable to adequately fund their own care and a bequest. For wealth levels sufficiently high, they are likely to fund their own care, as the value of public care becomes small compared to the value of private care and the total expenditures on LTC become small relative to their desired bequest level. This suggests that there will be an interior response that provides a measure of the equivalent dollar amount an individual assigns to receiving public care. Ultimately, to identify ψ_G , SSQ 4 asks respondents to provide the W that satisfies

$$\begin{aligned} & \frac{(\theta_{\text{ADL}})^{-\gamma}(\psi_G + \kappa_{\text{ADL}})^{1-\gamma}}{1 - \gamma} + \frac{(\theta_{\text{beq}})^{-\gamma}(W + \kappa_{\text{beq}})^{1-\gamma}}{1 - \gamma} = \\ & \frac{(\theta_{\text{ADL}})^{-\gamma}(z_1 + \kappa_{\text{ADL}})^{1-\gamma}}{1 - \gamma} + \frac{(\theta_{\text{beq}})^{-\gamma}(W - z_1 + \kappa_{\text{beq}})^{1-\gamma}}{1 - \gamma}, \end{aligned} \quad (7)$$

where z_1 is the optimal policy when there is no public care, as calculated in SSQ 3.

4. Identification

Because different utility functions are active in different states, different preference parameters control the marginal utility trade-offs that determine the decisions in each state of the world (see table 2 to see which parameters influence optimal decisions in each SSQ). For instance, SSQ 1 asks individuals to make a risky bet regarding consumption when healthy and explicitly rules out the potential that this decision could influence consumption in other health states. Since the relative risk aversion parameter γ is the only parameter that determines marginal utilities in the active states, this question identifies risk aversion. SSQ 2 examines the trade-off between having wealth when healthy and having it when in need of LTC for ADLs. This trade-off is optimally determined (abstracting from corner solutions for the moment) by equating marginal utility in the healthy state, as determined by γ , with marginal utility when in need of help, as determined by γ , θ_{ADL} , and κ_{ADL} . Utilizing the observed trade-offs at different wealth levels and state probabilities, we thus are able to identify the θ_{ADL} and κ_{ADL} necessary to align the model with SSQ responses. SSQ 3 examines a similar trade-off between wealth when in need of help with ADLs and wealth for a bequest, while SSQ 4 examines how the existence of a government LTC consumption floor effects the trade-off. In both of these questions, the model-implied optimal strategy is dictated by the marginal value of wealth in the ADL state (again, determined by γ , θ_{ADL} , and κ_{ADL}) and the marginal value of wealth allocated toward a final bequest (determined by γ , θ_{beq} , and κ_{beq}), while in the fourth SSQ the respondent must also take into account how the existence of a public care option affects this trade-off by determining how much he values public care (ψ_G).

C. Descriptive Analysis of SSQ Responses

In this section, we seek to describe how the SSQ responses will ultimately inform the formal estimation of preference parameters. We do so by analyzing the histograms of responses to the three variants of SSQ 2 and the three variants of SSQ 3, which are the questions most unique to our paper and most informative of the LTC and bequest utility parameters. Ameriks et al. (2018b) presents the histogram of responses for all SSQ questions.

In figure 8, we observe how individuals trade off consuming when healthy and consuming when in need of LTC, across wealth levels and different probabilities of needing care. As documented in figure 8A, when the probability of needing LTC is 0.25 and people had \$100K to allocate between the LTC state and the healthy state, most people allocated between \$30K and \$50K to the LTC state. Since the relative price of the LTC-state-contingent asset in the survey was set to $1/(1 - \pi)$, respondents

TABLE 2
LINK BETWEEN PARAMETERS AND SSQs

Question	Objective	Scenario Parameters	Preference Parameters
SSQ 1: lottery over spending	$\lambda^*: \frac{1}{1-\gamma}(W)^{1-\gamma} = \frac{.5}{1-\gamma}(2W)^{1-\gamma} + \frac{.5}{1-\gamma}[(1-\lambda^*)W]^{1-\gamma}$	(a) $W = \$100K$; (b) $W = \$50K$	γ
SSQ 2: allocation between ordinary and ADL states	$\max_{z_1, z_2} \pi \frac{z_1^{1-\gamma}}{1-\gamma} + (1-\pi) \frac{(\theta_{ADL})^{-\gamma} (z_2 + \kappa_{ADL})^{1-\gamma}}{1-\gamma}$	(a) $W = \$100K, \pi = .75$; (b) $W = \$100K, \pi = .50$; (c) $W = \$50K, \pi = .75$	$\gamma, \theta_{ADL}, \kappa_{ADL}$
SSQ 3: allocation between ADL and bequest states	$\max_{z_1, z_2} (\theta_{ADL})^{-\gamma} \frac{(z_1 + \kappa_{ADL})^{1-\gamma}}{1-\gamma} + (\theta_{BEQ})^{-\gamma} \frac{(z_2 + \kappa_{BEQ})^{1-\gamma}}{1-\gamma}$	(a) $W = \$100K$; (b) $W = \$150K$; (c) $W = \$200K$	$\gamma, \theta_{ADL}, \kappa_{ADL}, \theta_{BEQ}, \kappa_{BEQ}$
SSQ 4: indifference between public and private LTC	$W^*: (\theta_{ADL})^{-\gamma} \frac{(z_1 + \kappa_{ADL})^{1-\gamma}}{1-\gamma} + (\theta_{BEQ})^{-\gamma} \frac{(W^* - z_1 + \kappa_{BEQ})^{1-\gamma}}{1-\gamma} =$	Public care available	$\gamma, \theta_{ADL}, \kappa_{ADL}, \theta_{BEQ}, \kappa_{BEQ}, \psi_G$

NOTE.—The first column briefly summarizes the trade-offs, while the second lists the underlying optimization problem. The third column lists how question parameters were changed for different variations of each SSQ, where W is wealth and $1 - \pi$ is the probability of needing LTC. The z_1 in SSQ 4 is the optimal z_1 function calculated in SSQ 3. The fourth column lists the parameters that determine optimal responses in the model.




FIG. 8.—Allocations to the LTC state in SSQ 2 at different wealth (W) levels and probabilities of needing LTC ($1 - \pi$). Since the relative price of the LTC-state contingent asset in the survey was set to $1/(1 - \pi)$, respondents would have $1/(1 - \pi)$ times their chosen allocation available for spending in the LTC state.

would have $1/(1 - \pi)$ times their chosen allocation available for spending in the LTC state. Thus, these choices reflect the respondents' desire to have around \$120K–\$200K available when they need LTC and \$50K–\$70K available when they do not need LTC. Figure 8B shows the response distribution when the probability of needing LTC is increased to 0.5. In this case, more people allocate around \$40K–\$70K to the LTC state, resulting in around \$80K–\$140K when in need of care and \$40K–\$60K when healthy. Thus, when the probability of needing care is higher, people chose to sacrifice consumption when healthy to keep significant resources available when in need of care. For most respondents, both consumption when healthy and that when in need of care are lower, reflecting a tough trade-off on the margin. That is, respondents do not fix a set amount of wealth in either the healthy or the LTC state with the other state receiving the residual. Figure 8C plots the responses when people have only \$50K in wealth for the year. Most people choose to allocate between \$10K and \$30K to the LTC state, delivering \$40K–\$120K when in need of LTC and \$20K–\$40K when healthy. Again, even when

resources are tight, people are choosing to divert significant funds away from the healthy state toward the LTC state but are keeping substantial amounts for the healthy state as well. For all SSQ 2 variants, most respondents allocate resources so that they have more available in the LTC state than in the healthy state, often much more. These choices imply that LTC is a high-marginal utility state, generating low θ_{ADL} and negative κ_{ADL} in the estimation.

In figure 9, we observe how individuals trade off leaving money as a bequest and having wealth when in the ADL state, across wealth levels of \$100K (A), \$150K (B), and \$200K (C), the three variants of the bequest SSQ 3. The histograms show the amount that the individual would allocate toward the ADL state. Here, we clearly see that individuals react to the wealth level. Many respondents allocate almost all of their portfolio to the ADL state when wealth is \$100K. When wealth is \$150K, it is evident that the wealth constraint in scenario 3a severely restricts spending in the ADL state, as evidenced by the large mass of individuals responding with allocations to the ADL state above \$100K. Similar patterns repeat when wealth is \$200K, with significant response mass above \$150K. Furthermore, even at \$200K, many individuals gave zero bequest. This finding suggests that many people view LTC as a primary reason to save late in life




FIG. 9.—Allocations to the LTC state in SSQ 3 at different wealth (W) levels.

at lower wealth levels, with bequest motives becoming more important at higher wealth levels. These responses are consistent with the view that bequests are considered a luxury good while LTC is a necessary good, which should be reflected in estimates of κ_{ADL} and κ_{beq} .¹⁵ It is harder to read directly how these responses translate into θ_{ADL} and θ_{beq} , especially as parameters are jointly estimated using all SSQs. These figures indicate, however, that estimators that target SSQ moments are likely to result in parameter estimates that indicate a strong desire to spend when in need of LTC, compared to leaving a bequest. Having provided some descriptive evidence on how the SSQ responses will inform parameter estimates, we turn to the formal estimation strategy.

V. Estimation Methodology

We develop a two-stage MSM (method of simulated moments) estimator that is similar to those used in Gourinchas and Parker (2002), Laibson, Repetto, and Tobacman (2007), De Nardi, French, and Jones (2010), French and Jones (2011), and Lockwood (2018) to estimate parameter set Γ . The set $\Gamma := [\Xi, \Theta]$ is divided into two subsets, with the first subset, Ξ , consisting of parameters externally estimated without the use of the structural model (e.g., income, health transitions, and health costs) and the second parameter subset, $\Theta := \{\gamma, \theta_{\text{ADL}}, \kappa_{\text{ADL}}, \theta_{\text{beq}}, \kappa_{\text{beq}}, \psi_G\}$, consisting of preference parameters that are estimated using moments generated by simulating the structural model.

First, we set some parameters to values informed by the literature and explore robustness to these externally chosen values. We restrict $\beta = 0.97$ and $\omega_G = 30\text{K}$, as the empirical strategy using SSQs was not designed to estimate these parameters. The VRI sample is wealthy enough that almost no one chooses to go on government welfare when healthy, and thus parameter estimates are identical for a wide range of ω_G .¹⁵

A. Estimated Inputs for the Model

1. Income

Income profiles are estimated from VRI Survey 1. In Survey 1, respondents report their income flows as the sum of labor income, pension and disability payments, and social security payments. For each age, we

¹⁵ Parameter β is not identified from the SSQs. If β is estimated when targeting wealth moments, it is usually estimated to be close to 1 in order to help the model generate the large degree of wealth present in the upper deciles of the empirical wealth distribution. The effect of this on the other estimated parameters is not large. Robustness to alternative values of β is presented in sec. IX. We preset $\omega_G = 30\text{K}$ for all exercises because that is the point estimate—although with very large standard errors—when including this parameter to be estimated in our preferred baseline.

assign respondents to an income quintile based on their current rank among individuals of the same age. Using this cross section of income, we use a quintile regression to estimate the age profile of earnings as a polynomial of age and sex. Since life-cycle parameters are estimated off the cross section, we cannot control for cohort effects. The model groups consumers into five income groups with deterministic age-income profiles determined by the estimated coefficients. This allows us to capture income changes during retirement but abstract from income fluctuations as a source of uncertainty. The estimated age profiles of income for each quintile are presented in the appendix (sec. A1).

2. Health Transitions

Health transitions are estimated using HRS waves 2–10. We first apply the sample selection criteria of the VRI to the HRS data to obtain a population with similar observable characteristics. This accounts for the fact that the VRI is wealthier and more educated than the US population and that health-transition probabilities vary substantially in these dimensions. Then we construct the defined health states ($s = 0, 1, 2$) from two sets of questions. The first utilizes self-reported subjective health-status questions to classify individuals into good or bad health ($s = 0$ or $s = 1$).¹⁶ The second set of questions is used to determine whether an individual is in the LTC/ADL state ($s = 2$). This set of questions presents five ADLs and asks whether respondents receive help with any of the five activities. If the respondent answers yes to any of these questions, then we define that respondent to be in the ADL health state.¹⁷ Transitions are then estimated using a maximum likelihood estimator, with more information on the estimation methodology and the resulting estimates provided in the appendix (sec. A2).

Given the HRS data, there are two alternative LTC/ADL state definitions that are feasible: needing help with ADLs and being in a nursing home. We prefer to use the “receive-help” variable to define the ADL health state, for two reasons. First, being in a nursing home is too restrictive to represent the general health status of needing help with ADLs. There are many people who need help with ADLs who do not reside in a nursing home, and even when in need of help with ADLs, people

¹⁶ Individuals are defined as in good health if they report health being good, very good, or excellent and are defined to be in bad health if they report health being poor or fair. Self-reported assessments of health status have strong predictive power for future health realizations.

¹⁷ Results on saving behavior are robust to using receiving help with at least two or at least three ADLs as the cutoff for being in the ($s = 2$) state; the effect of a lower probability of needing help with ADLs and the increased mortality rate associated with the more severe definition of $s = 2$ is offset by the increase in persistence of the $s = 2$ state. Results of these alternative thresholds are presented in sec. IX.

may spend on goods other than LTC. Second, it is difficult to interpret a report of needing help but not receiving help as really needing help. Thus, we choose the more general needing-help-with-ADLs status but the more stringent requirement that they actually report receiving help. For consistency, this is the same ADL definition presented in the VRI survey.

3. Health Expense

The health expense distributions are estimated using the 2010 HRS data. We model the mandatory out-of-pocket health expenditures as being distributed lognormally and estimate the mean and standard deviation of the lognormal distribution conditional on age, sex, and health state. Because idiosyncratic health cost is independently and identically distributed conditional on states, a single year of cross-sectional data is sufficient for estimation. For the VRI sample studied here, out-of-pocket (post-insurance) health cost shocks are relatively small, compared to total wealth, and thus do not substantively drive saving in the model. We use a measure of health costs that includes many subcategories of expenditure, and it is difficult to construct an ideal measure of exogenous expenditure for our purposes from the HRS; reassuringly, the key results in this paper are virtually unchanged over a large range in assumptions about these medical expenditure shocks, as documented in section IX. For more information on the estimation of costs and the resulting estimates, see the appendix (sec. A2).

B. Second-Stage Estimates

In the second stage, we apply the MSM estimation procedure. The moments we use to estimate the model are derived from two distinct survey measurements. The first set of moment conditions is derived from behavioral data. We target age-conditional wealth percentiles, which are frequently used to estimate similar life-cycle saving models. A second set of moments is derived from SSQ responses. We first describe the estimator and then detail the construction of both sets of moments and present the comprehensive moment set we target in our baseline estimation. In section VI, we present and analyze the resulting baseline parameters. Furthermore, by design of the SSQs, it is possible to estimate the model using only wealth or only SSQ moments, which we explore in section VIII.A.

We define $X = (X_i)_{i=1}^I$ as the collection of measurements for all individuals, including behavioral responses, SSQ responses, and state variables; $x_i \subseteq X_i$ is used to denote relevant subsets of individual i 's data. The moments used for estimation are the difference between statistics generated

by the structural model $(m(\tilde{\mathbf{Z}}, \Theta, X))$ and empirical data $(s(X))$, defined as $g(\tilde{\mathbf{Z}}, \Theta, X) = \mathbb{E}[m(\tilde{\mathbf{Z}}, \Theta, X) - s(X)]$. We estimate second-stage parameters $\hat{\Theta}$ that minimize a GMM quadratic objective function with moments $g(\tilde{\mathbf{Z}}, \Theta, X)$. That is, the second-stage estimator with weighting matrix W is

$$\hat{\Theta} = \arg \min_{\Theta} g(\tilde{\mathbf{Z}}, \Theta, X)' W g(\tilde{\mathbf{Z}}, \Theta, X). \quad (8)$$

To estimate the model with the optimal weighting matrix, we use the standard two-step feasible MSM approach. In the first step, we minimize the objective function defined in equation (8), using the identity weighting matrix, and denote the minimizing parameter set $\hat{\Theta}_1$. Using this parameter set, we calculate the moment vector $g(\tilde{\mathbf{Z}}, \hat{\Theta}_1, X)$ and the implied first-step covariance matrix $\hat{\Omega}_1$. We use the inverse of the first-step covariance matrix as the second-step optimal weighting matrix $\hat{W} = \hat{\Omega}_1^{-1}$. We then minimize according to equation (8), using \hat{W} to estimate the final parameter set $\hat{\Theta}$.

1. Wealth Moments

As is common with many life-cycle studies of late-in-life savings (e.g., Gustman and Steinmeier 1986; De Nardi, French, and Jones 2010; Lockwood 2018), the first moment set consists of asset percentiles, conditional on a set of state variables x . The simulated wealth percentile p conditional on state variables x is denoted as $\alpha_x^p(\mathbf{Z}, \Theta, X)$, while a_i denotes individual i 's empirical wealth holdings. The wealth moments conditional on x are

$$g_{p,x}(\tilde{\mathbf{Z}}, \Theta, X) = \mathbb{E} \left[\mathbb{I}_{\{a_i < \alpha_x^p(\mathbf{Z}, \Theta, X)\}} - p \mid x_i = x \right], \quad (9)$$

an expression that can easily be converted to an unconditional expectation through the law of iterated expectations.

For our baseline, we define the moment conditions as the 25th, 50th, and 75th percentiles conditional on age t (a_t^p). In practice, we aggregate the age profiles into disjoint 3-year intervals, so that $t \in \{55-57, 58-60, \dots, 88-90\}$.¹⁸ This aggregation is done to smooth noise in the empirical asset profile that we observe in the cross-sectional data. The wealth moments that we construct from cross-sectional data are subject to mortality bias, since wealthier people tend to live longer. Chetty et al. (2016) show that mortality rates are much lower at these ages for high-income individuals. Thus, mortality bias is likely to be less severe in the VRI sample than in the

¹⁸ We stop matching wealth moments when the sample for the age group becomes too small, defined as less than 20 individuals. This first occurs in the 91-94 age bin. Results are not sensitive around a reasonable range of the chosen cutoff.

general US population.¹⁹ Given the three targeted wealth percentiles for each of the 12 age intervals, the wealth moment set has 36 moments.

To generate the simulated wealth moments, we first solve the model to compute optimal decision rules as a function of all relevant state variables, given parameters. We simulate the model for a large number of individuals ($N = 10,000$) and then compute moments, using the simulated behavior. To set the initial conditions for the simulation, we sample (with replacement) N individuals from the VRI data X , in which an individual is characterized by his or her vector of idiosyncratic state variables. Then, for each individual, we draw relevant shocks from the $\hat{\mathbf{Z}}$ parameterized stochastic processes and simulate the behavior implied by the computed optimal policies (given parameters Θ). We then aggregate these individual behaviors to construct the simulated population moments $m(\hat{\mathbf{Z}}, \Theta, X)$, pooling observations from different calendar years into age bins.

2. SSQ Moments

The second moment set is constructed entirely from SSQ responses. The structural model has six preference parameters to be estimated: the relative risk aversion parameter γ ; LTC-state utility parameters θ_{ADL} and κ_{ADL} ; bequest utility parameters θ_{beq} and κ_{beq} ; and public care aversion, ψ_{G} . Using responses to the four types of SSQs described in section IV.B and their iterations at different wealth levels and state probabilities, we are able to identify these six parameters. Since the survey poses two variants of SSQ 1, three variants of SSQ 2, three variants of SSQ 3, and one variant of SSQ 4 and we are estimating six preference parameters, using SSQ data alone we have an overidentified system suitable for estimation of these six preference parameters.

As moments, we match the model-implied allocations to the empirical mean of each of the nine SSQ variants, indexed by $m \in \{1, 2, \dots, 9\}$. Let z_m^i be individual i 's response to SSQ variant m . The theoretical counterparts, the optimal SSQ answers conditional on preference parameters, are denoted $s_m(\Theta)$. We then write each moment as

$$g_m(\hat{\mathbf{Z}}, \Theta, X) = \mathbb{E}[s_m(\Theta) - z_m^i], \quad (10)$$

providing us with nine SSQ moments.

3. Combining Wealth and SSQ Moments

In the previous two sections, we described the specification of wealth and SSQ moments. In this section, we describe our baseline estimation

¹⁹ Recall that the estimated health-transition probabilities in the model are based on VRI-eligible HRS respondents, so they take into account the expected mortality experience of the VRI population.

procedure, which uses both SSQ and wealth data by combining these moments into a single moment vector that will be used to estimate the model. Utilizing both sources of information disciplines the estimator to match wealth data and SSQ data, with each source of data likely containing unique information. To combine the two sources of data, we concatenate the wealth moments formed from matching the cross-sectional wealth distribution percentiles and the SSQ moments formed by matching the empirical SSQ responses, resulting in a total of 45 moment conditions.

Letting $g_v(\tilde{Z}, \theta, X)$ denote the set of moments constructed from wealth data and $g_s(\tilde{Z}, \theta, X)$ denote a set of moments constructed from SSQ data, the baseline joint estimation's moment set is

$$g_J(\tilde{Z}, \theta, X) = \begin{bmatrix} g_v(\tilde{Z}, \theta, X) \\ g_s(\tilde{Z}, \theta, X) \end{bmatrix}. \quad (11)$$

By design, SSQ moments provide strong identification of the preference parameters, and thus the optimal weighting matrix assigns a disproportionate amount of weight to the SSQ moments, resulting in tight match of SSQ responses but a slight underprediction of wealth accumulation. In order to better match these wealth moments, which are the common target of other studies, and to help facilitate comparison to the literature, we make use of a parameter, λ , that allows us to control the relative importance we assign to each of the moment sets $g_v(\tilde{Z}, \hat{\theta}, X)$ and $g_s(\tilde{Z}, \hat{\theta}, X)$. By overweighting the wealth moments relative to the statistically optimal weighting matrix, we arrive at our baseline estimator.²⁰

VI. Parameter Estimates

A. Model Fit

Before presenting the parameter estimates, we show the fit of the model to the data. Figure 10 documents the model fit from our baseline estimation that jointly targets both wealth (*A*) and SSQ (*B*) moments. For exposition, we present the SSQ moments on a $[0, 1]$ scale, by normalizing the mean response by the maximum possible response.²¹ Overall, the fit is good: the percentiles of the wealth distribution are matched well, as are the SSQ moments. We first analyze the baseline estimates and then analyze the distinct contributions of the SSQs and wealth moments in section VIII.

²⁰ See the online appendix for results using the optimal weighting matrix. Baseline estimates are quantitatively similar to those using the optimal weighting matrix. This similarity holds for a wide range of λ , suggesting that the implications of preferences for saving, expenditure, and bequests are robust to this choice.

²¹ Because there is no maximum response to SSQ 4, we set W_m to be the 95th percentile of the response distribution.




FIG. 10.—Model fit when jointly targeting wealth and SSQ moments; $m_p = m$ th percentile.

B. Baseline Estimated Parameter Values

Table 3 presents the parameter estimates for our baseline case that combines the wealth and SSQ moments. Of particular interest in the estimation are four groups of parameters. First, the estimate of risk aversion (γ) is of independent interest but is also important because it is jointly estimated and strongly affects the estimated values of other preference parameters. Second, the estimates of health-state utility (θ_{ADL} and κ_{ADL}) presented are the first we are aware of to estimate this functional form applied specifically to the LTC health state. Third are the bequest parameters θ_{beq} and κ_{beq} . Finally, ψ_G controls the degree to which there is public care aversion.

The coefficient of relative risk aversion is estimated to be 5.27. This value is somewhat larger than that typically used in the literature and somewhat smaller than that typically estimated using similar survey techniques. For example, in an exercise using a similar model and very different data, De Nardi, French, and Jones (2010) estimate $\gamma = 3.8$, while Kimball, Sahm, and Shapiro (2008) estimate a mean γ of 8.2 when using responses in the HRS to hypothetical lotteries over future income streams.

In examining the estimated preferences when targeting the joint moments, it is striking that the health-state utility function implies very different marginal utility of wealth in the state of the world when an individual needs LTC and when LTC is not needed. The estimated $\kappa_{\text{ADL}} = -37.44$ is negative and large, suggesting that LTC is viewed as a necessary good with a spending floor of about \$37K.²² The estimated $\theta_{\text{ADL}} = 0.67$ implies a high marginal utility of expenditure in the LTC state, especially when combined with κ_{ADL} . Note that if $\kappa_{\text{ADL}} = 0$, relative to ordinary consumption, $\theta_{\text{ADL}} > 1$ would imply that LTC expenditures provide less marginal utility for each dollar spent, while $\theta_{\text{ADL}} < 1$ would imply a higher marginal utility for each dollar of expenditure. Individuals with $\theta_{\text{ADL}} > 1$ and $\kappa_{\text{ADL}} < 0$ would view expenditure when in need of LTC as a strong necessity and so would optimally desire to consume the necessary amount, but not much more. Thus, the effect of precautionary savings on wealth accumulation would be similar to medical expenses such as those presented in De Nardi, French, and Jones (2010). For individuals with $\theta_{\text{ADL}} < 1$, marginal expenditures in the LTC state would have a high value, so there is a motive to have additional spending when in need of care even at expenditures above the necessary amount.

For a given amount of spending, individuals have a higher marginal utility of spending in the LTC state than from leaving a bequest. First,

²² In our baseline, we set the minimum private LTC expenditure $\chi = 40,000$. The estimation results are independent of the choice of χ for any $\chi \leq -\kappa_{\text{ADL}}$, and χ is quantitatively irrelevant when greater than but close to κ_{ADL} , as is the case here.

TABLE 3
JOINT PARAMETER ESTIMATES: BASELINE MODEL

γ	θ_{ADL}	κ_{ADL}	θ_{beq}	κ_{beq}	ψ_G
5.27 (.05)	.67 (.37)	-37.44 (.22)	1.09 (.25)	7.83 (.48)	77.43 (9.49)

NOTE.—Standard errors are reported in parentheses.

$\kappa_{\text{beq}} = 7.83$, which implies that bequests are viewed as a luxury good. Together with a positive κ_{beq} , the estimate of $\theta_{\text{beq}} = 1.09$ implies a low marginal utility of bequests, relative to LTC. As discussed in section VIII.B, estimates in the literature, like those from Lockwood (2018) and Koijen, Van Nieuwerburgh, and Yogo (2016) find stronger bequest motives.

Because expenditure when in need of LTC is so highly valued, $\psi_G = 77,430$ is estimated to be large. To facilitate comparison to the literature, this estimate implies a utility level equivalent to that provided by a \$24,395 government-provided public care expenditure in a model with the same risk aversion but without state-dependent preferences.²³

VII. Behavioral Implications
of Estimated Preferences

To help further interpret the parameter values in economically meaningful ways and to investigate the contribution of the different saving motives in determining life-cycle wealth patterns, in this section we use the model to document behavior induced by the estimated preferences. We first do so in a simple static optimization problem to give a sense of the curvature of the estimated utility functions, before turning to an examination of expenditure in the estimated full model.

A. An Illustrative Synthetic Static Choice Problem

As a first pass on interpreting the estimated parameters, we plot optimal allocations in a simple synthetic static choice problem, in which an individual regards consumption when healthy, consumption when in need of LTC, and bequests as three different “goods” that can be purchased simultaneously at a relative price of 1 and are valued according to the corresponding estimated utility functions. This problem is not one that maps to a situation ever faced by individuals; for example, there is no probability weighting to reflect the likelihood of being in each state. It is nonetheless a convenient illustration device to present the marginal

²³ To calculate this expenditure equivalent in a model without the health-state utility function, we find the expenditure level ψ that would equate utility across the two specifications: $\psi^{1-\gamma}/(1-\gamma) = (\theta_{\text{ADL}})^{-\gamma}[(\psi_G + \kappa_{\text{ADL}})^{1-\gamma}/(1-\gamma)]$.

utility of expenditures by health status implied by the estimated parameters before turning to the full model that incorporates dynamics, risk, and government insurance. Specifically, the static allocation problem for a person with wealth W is

$$\begin{aligned} \max_{z_1, z_2, z_3} \quad & \frac{(z_1)^{1-\gamma}}{1-\gamma} + \frac{(\theta_{\text{ADL}})^{-\gamma}(z_2 + \kappa_{\text{ADL}})^{1-\gamma}}{1-\gamma} + \frac{(\theta_{\text{beq}})^{-\gamma}(z_3 + \kappa_{\text{beq}})^{1-\gamma}}{1-\gamma}, \\ \text{subject to } & z_1 + z_2 + z_3 \leq W, \\ & z_1, z_2, z_3 \geq 0, z_2 \geq -\kappa_{\text{ADL}}, \text{ and } z_3 \geq -\kappa_{\text{beq}}. \end{aligned} \quad (12)$$

Recall that this synthetic choice model is meant to summarize the parameters' implications succinctly, not to provide quantitative predictions reflecting an actual choice problem. The main advantage of this simple choice problem is transparency, as the interaction between motives that is expressed in the full dynamic model makes it challenging to easily determine the relative strength of each motive.

Figure 11 plots the resulting optimal allocations. The most striking feature of this figure is that ADL-state expenditure is allocated the majority of wealth up to rather high wealth levels. Because ADL-state expenditure




FIG. 11.—Expenditure in synthetic static choice problem.

is a necessary good, at least $-\kappa_{\text{ADL}}$ dollars are always allocated to this motive, giving rise to an initial ADL allocation share nearly equal to 1 that falls as wealth is allocated to other expenditures to take advantage of diminishing marginal utility from any one type of expenditure. In addition, we see that because bequests are estimated to be a luxury good, there is initially no spending on bequests until wealth is sufficiently high that the bequest motive becomes active. The effect on expenditure of the nonhomotheticity parameter κ diminishes at higher wealth levels, and the differences in θ become more dominant. Even at high wealth levels, people spend the most on the ADL-state good, then the ordinary-state good, followed by a bequest. Although the outcome of a simple synthetic static optimization problem, these results highlight that the estimated utility functions reflect a strong desire to spend when in need of help with LTC and a relatively weak preference for leaving a bequest, which will contribute to life-cycle saving behavior in the full dynamic model with health and mortality risk.

B. Implied Expenditures Using the Baseline Model

In this section, using the full estimated model, we analyze expenditure behavior over the life cycle and in the cross section.²⁴ The simulated expenditures show the implications of the parameter estimates on an important nontargeted dimension and provide credibility for the quantitative behavior of the model.

We begin by documenting the cross section of expenditure across health status in figure 12. The heterogeneity in expenditure comes from differences in income, wealth, age, sex, and the realization of health shocks. When in good health, the modal expenditure is around \$40K, with a long upper tail. As alluded to in the analysis of the estimated parameters, spending is higher when in need of LTC than when healthy. When in need of LTC, over 60% of people spend around \$80K–\$100K. Furthermore, there are around 15% of people spending over \$150K when in need of LTC and 5% spending over \$200K. Median spending when in need of LTC increases by about \$50K per year, from around \$40K when in good health to \$90K.

We model total expenditure in the ADL state, which includes food, shelter, and so on, as well as spending on ADL-related goods and services. The measure reflects potentially higher costs of ordinary goods in the ADL state as well as changes to spending that arise from changing survival probabilities that are associated with the health state. Nonetheless, to get a sense of comparison to the data, we can compare to the cost

²⁴ Expenditure is defined as $c + h$, summing all voluntary expenditure and the mandatory out-of-pocket health cost shocks. These expenditures result from the same simulations used to construct the wealth moments.




FIG. 12.—Distribution of expenditure. This figure plots the cross sections of expenditure for those in good health and for those in need of help with ADLs.

of LTC as reported by Genworth (2016). There is a wide range of nursing home costs: an annual stay in a private room in a nursing home averages \$92K and ranges between \$60K and \$250K, reflecting varying costs and qualities. Not all individuals who need help with LTC move to nursing homes. Those with less severe needs may purchase formal in-home care, which typically costs about \$50K per year. Thus, the change in spending when in need of LTC predicted by the model aligns well with the costs associated with obtaining extra care. Beyond spending on formal LTC, there are various ways that needing help with ADLs can lead to increased spending. For example, people may purchase more expensive prepared food if cooking is difficult, they may hire help to clean the house and care for landscaping, and/or they may install devices around the house to help with mobility. The increase in spending in our model when people are in need of help with ADLs is designed to capture costs directly related to LTC as well as these other reasons that living with ADL impairments is expensive.²⁵

Figure 13 presents the distribution of expenditure by health status conditional on age. The distribution of expenditure for people in good health is nearly constant across age. The 90th percentile is around \$100K per year, while the 10th percentile is around \$30K. Expenditure when in need of LTC is also fairly constant across age, with substantially higher skewness at all ages and more variation in the upper percentiles. These expenditure patterns are consistent with the interpretation that spending when in need of help with ADLs is viewed as a necessary good with an additional intensive-margin valuation.

Figure 14 puts this all together by plotting the share of total expenditure for each age by ADL status. This calculation combines the fraction of people with each health status and the distribution of expenditures by health status. The most dramatic change over the life cycle is the decrease in the share of expenditure by those in good health and the increase in the share associated with those needing help with ADLs at older ages. At younger ages, very few people need help with ADLs, and it is not until around age 80 that those people account for over 10% of total

²⁵ Since we model spending when in need of LTC and not spending on LTC, there is no direct mapping from consumption and medical spending in our model to health-related out-of-pocket expenditures in surveys like the HRS or Medicare Current Beneficiary Survey (MCBS), which limits our ability to compare model predictions to data. Future advances in measuring total consumption and expenditure with large samples in the appropriate age range would help, but current consumption data in the United States are limited and not available in sufficiently large samples to provide good measures of state-contingent spending. Future research could systematically examine the multiple sources of changes in total expenditure by health state: direct spending on ADL-related costs, different prices paid for ordinary spending, a change in the composition of spending, and a change in expected lifetime.

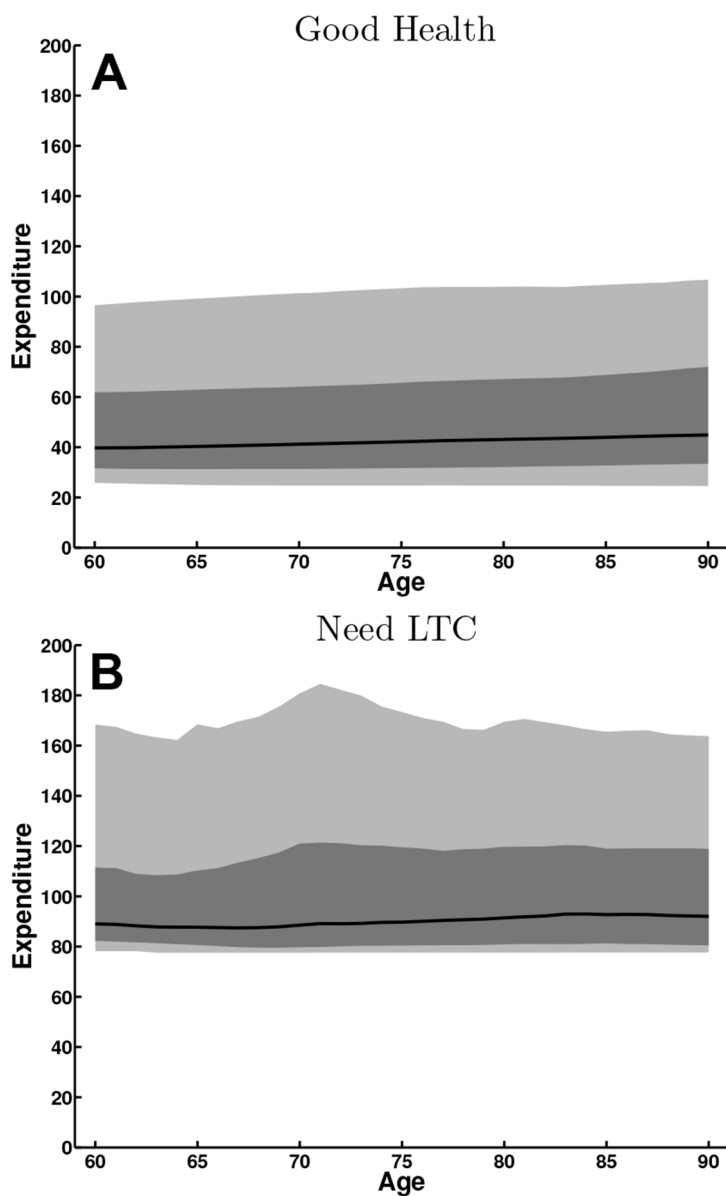


FIG. 13.—Distribution of expenditure by age. This figure plots expenditure by age for those in good health and for those in need of help with ADLs. The solid line represents median expenditure, and the shaded areas represent the 25th–75th (darker) and the 10th–90th (lighter) percentiles.




FIG. 14.—Share of expenditure by age and health status.

expenditure. After age 80, however, the expenditure share associated with those needing LTC increases rapidly. The probability of needing care increases at these ages, and the expenditure by those needing care is significantly higher than that by those in good health. At age 90, those needing help with ADLs account for about one-third of expenditures. This suggests that risks and associated preferences combine to make LTC a significant driver of saving for the elderly. Taken together, these figures show that the estimated model predicts reasonable expenditure patterns in the cross section and over the life cycle, at the individual and population levels.

Finally, to document the relative strength of the different saving motives, we present in table 4 percentiles of the distribution of bequests under alternative parameter values. Bequests are one useful variable to describe saving over the life cycle, since they are the result of a lifetime of asset accumulation. Similar patterns are present in other measures of life-cycle saving, for example, maximum wealth. The differences presented in table 4 reflect the relative strength of the LTC and bequest utility functions estimated with the wealth and SSQ data using the full model, as opposed to the synthetic static choice problem featured above. The

TABLE 4
DECOMPOSING SAVING MOTIVES: BEQUESTS UNDER ALTERNATIVE PREFERENCES

	MEAN	PERCENTILE				
		10th	25th	50th	75th	90th
Baseline	505	0	157	455	707	1,024
No bequest utility	336	0	0	158	530	891
No health-state utility	302	0	59	200	408	677
No bequest or health-state utility	235	0	0	74	340	638

NOTE.—This table presents percentiles of the bequest distribution in thousands of dollars under different parameterizations of health-state-dependent and bequest utility functions.

first row shows bequests under the baseline parameterization, the second row shows bequests when there is zero utility from leaving a bequest ($\theta_{\text{beq}} = \infty$), the third row shows bequests when utility when in need of help with ADLs is equal to utility when healthy ($\theta_{\text{ADL}} = 1, \kappa_{\text{ADL}} = 0$), and the fourth row combines these alternatives by showing bequests when there is no bequest utility and no health-state-dependent utility ($\theta_{\text{beq}} = \infty, \theta_{\text{ADL}} = 1, \kappa_{\text{ADL}} = 0$).

Comparing these parameterizations provides insight on the contribution of the various saving motives to overall late-in-life wealth accumulation. In the baseline model, realized bequests reflect an active desire to leave a bequest and an incidental bequest driven by precautionary saving combined with mortality risk. These motives interact to reinforce each other: with strong bequest motives, the cost of ex post oversaving for the ADL health state is lower, and strong ADL-state utility induces high saving rates that may result in a large incidental bequest. Because of this interaction, this exercise does not provide a strict decomposition. Nonetheless, the magnitude of the differences is informative of the strength of the motives. These interdependencies highlight the benefit of SSQs that can help to separately identify preferences in light of these interactions.

When the bequest motive is shut off, the average bequest is two-thirds its baseline level, suggesting that the majority of saving over the life cycle is driven by precautionary motives related to health and death and that most bequests are incidental, as opposed to being driven by an active warm-glow bequest margin. Furthermore, average bequests are larger in the model with no bequest motive than in the model with no health-state-dependent utility. In the cases with zero bequest utility, bequests are, on average, 43% larger with the LTC utility function than without. On average, the model without either the bequest motive or the LTC utility function generates just under half of the baseline bequest, indicating that longevity risk is a significant contributor to realized bequests but not the whole story quantitatively. There are also different responses to the alternative parameterizations across the bequest distribution. The median bequest is lower in the case in which there is no bequest utility than

when there is no health-state-dependent utility, while the 75th and 90th percentiles of the bequest distribution are larger. This pattern arises because the means-tested public care option dissuades lower-wealth individuals from saving for LTC reasons and because health-state utility encourages rapid dissaving when in need of care toward the end of life that is subject to diminishing within-period returns. Overall, table 4 shows that both precautionary motives and bequests contribute substantially to wealth accumulation over the life cycle and that health-state-dependent utility related to LTC needs is a quantitatively significant determinant of saving patterns at older ages.

C. *Fit of Nontargeted Wealth Moments*

To demonstrate the saving behavior across the wealth distribution implied by the baseline estimates, figure 15 plots the model fit to nontargeted wealth moments. Even though the 10th percentile was not targeted in the estimation, the model fit to data is almost exact. Preference parameters estimated by targeting the higher wealth moments successfully replicate savings patterns of the less affluent. The model does a good job of




FIG. 15.—The 90th-, 75th-, 50th-, 25th-, and 10th-percentile wealth moments in the model (dotted lines) and the data (solid lines).

“out-of-sample” prediction for low levels of wealth, which corresponds to a large fraction of the US population who may be the target of policy changes and innovations to promote retirement saving. The use of SSQs permits predictions over households who do not—because of their level of financial resources or institutional and market constraints—face the same economic trade-offs as members of the VRI sample. We explore this further by analyzing model predictions for the HRS sample in section IX.B. In contrast, the model predicts less wealth than appears in the data at the 90th percentile, typically underpredicting wealth by around 10%–15%. Our interpretation is that the model is missing certain features that are particularly important for the wealthiest members of the VRI sample, which itself oversamples high-wealth individuals relative to the US population.

VIII. Comparison to Alternative Parameter Estimates

A. *Matching Only SSQ or Only Wealth Moments*

In our baseline estimation, we jointly target the wealth moments that are traditionally used in this literature and our newly created SSQ moments. In order to disentangle the relative contributions of the SSQs in the estimation results, in this section we present and analyze the estimated preferences that result from matching wealth moments or SSQ moments exclusively. We are able to use the SSQ moments exclusively because we have collected sufficient nonbehavioral data to identify the preference parameters solely from the SSQ responses. Although our preferred estimates use both wealth and SSQ data, we do not need to augment the nonbehavioral data with observed behaviors to gain identification. We are unaware of any existing study that has undertaken a similar estimation of a structural life-cycle model without requiring behavioral data for parameter identification. The SSQ data serve two roles. First, they corroborate the importance of LTC that the wealth data alone suggest. Second, they substantially improve the precision of the preference parameter estimates, reducing standard errors to levels that permit concrete, instead of suggestive, conclusions.

Although it is difficult to compare across estimates that target different moments, some broad messages are clear. Every set of estimates suggests that the marginal utility of expenditures when in need of LTC is larger than that from bequests for all relevant wealth levels. The parameter κ_{ADL} is always estimated to be negative and large in magnitude. Often, θ_{ADL} is smaller than θ_{beq} , and in the SSQ-moment case when they are almost equal, the large negative κ_{ADL} dominates the positive and small κ_{beq} . Thus, it is not just because we target SSQs that we find the importance of the precautionary saving motive induced by LTC risk. Hence, many features of

the data strongly support allowing for a nonhomothetic LTC health-state utility function, that spending when in need of LTC is viewed as a necessity, and that such spending is highly valued on the margin. Moreover, direct analysis of the distribution of SSQs responses implies that the preference for leaving a bequest is not as strong as the wealth data alone suggest.

1. Model Fit

First, consider targeting only the wealth moments. Figure 16A presents the model-generated and empirical 25th, 50th, and 75th percentiles of the wealth distribution across ages. The model-generated data match the cross-sectional wealth distribution well for most ages. As displayed in figure 16B, the parameters that best match the wealth moments generate SSQ moments that are somewhat distant from those measured in the data. Roughly speaking, the resulting parameters suggest individuals who are more risk averse and have stronger desires to spend when in need of LTC and to leave bequests than the SSQ data imply.

Now consider targeting only the SSQ moments. As can be seen in figure 17, the SSQ moments are hit almost exactly. The success in fitting these moments should not be surprising, given that SSQs were designed to ensure identification, though the SSQ responses were not guaranteed to be mutually consistent. When only the SSQ moments are targeted, the fit to the wealth moments deteriorates, with the model predicting undersaving for the 75th percentile and oversaving for the 50th percentile and, especially, the 25th percentile.

A model's ability to match wealth moments does not necessarily mean that preferences and all associated saving motives are well identified. Including SSQ moments introduces more information on saving motives and disciplines the channels through which a model can match wealth moments. The tension in matching both sets of moments simultaneously highlights the need for further model development. We have started here by introducing the nonhomothetic health-state utility function. The information contained in the SSQs provides a source of identification for even richer models.

2. Estimated Parameter Values from Alternative Moments

Table 5 documents the parameter estimates that result from exclusively targeting wealth or SSQ moments. Because of the nonlinearity of the model, the joint estimates are neither a convex combination nor contained within the interval of the wealth-only or SSQ-only estimates.

Compared to the baseline, when the SSQ moments alone are matched the estimation procedure delivers a lower γ and higher θ_{ADL} and θ_{beq} . The

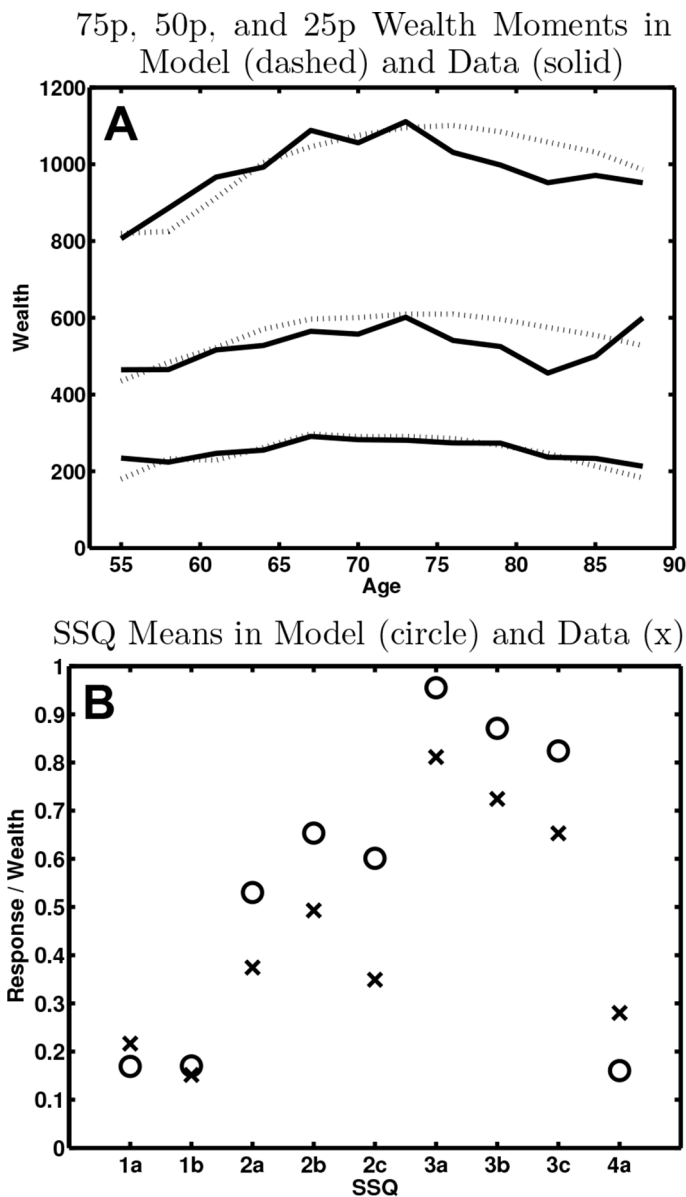


FIG. 16.—Model fit when exclusively targeting wealth moments; $np = n$ th percentile.

estimates targeting wealth moments result in a higher risk aversion coefficient and a higher marginal utility when in need of LTC, generated by both lower θ_{ADL} and more negative κ_{ADL} . In the wealth estimation, we find that most parameters are estimated with very large standard errors,

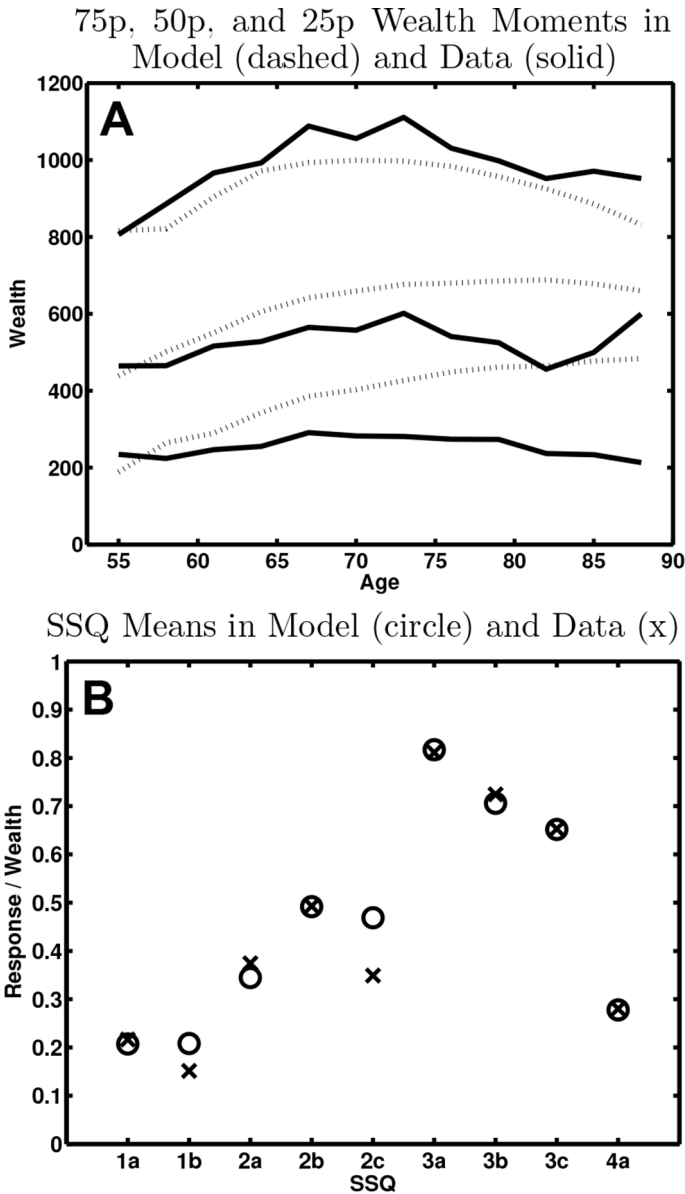


FIG. 17.—Model fit when exclusively targeting SSQ moments; $np = n$ th percentile.

which reflects the fundamental difficulty in identifying richer preference parameters from wealth data alone. This difficulty was discussed extensively in Ameriks et al. (2011) and commented on by De Nardi, French, and Jones (2010) in discussion of their table 2. Additionally, because we have only a

TABLE 5
ESTIMATED PARAMETERS: ALTERNATIVE ESTIMATES

	γ	θ_{ADL}	κ_{ADL}	θ_{beq}	κ_{beq}	ψ_G
Joint estimation: baseline model	5.27 (.05)	.67 (.37)	-37.44 (.22)	1.09 (.25)	7.83 (.48)	77.43 (9.49)
Wealth moments only	4.41 (3.85)	.51 (8.86)	-53.12 (4.75)	1.17 (1.07)	14.06 (7.26)	79.71 (7.18)
SSQs only	3.56 (.03)	1.77 (.18)	-63.97 (1.32)	1.61 (.14)	.34 (1.73)	91.68 (2.61)

NOTE.—This table presents parameter estimates for the estimation targeting jointly both sets of moments, targeting wealth moments only, and targeting SSQ moments only. Standard errors are reported in parentheses.

cross section, we cannot control for cohort effects. The use of SSQs allows us to design questions that overcome this limitation.

Comparing parameters estimated to match alternative moments.—It is difficult to compare the economic interpretation of estimated parameters in isolation because the value of γ affects the interpretation of the θ and κ parameters, just as the value of κ_{ADL} affects the interpretation of the value of θ_{ADL} . To compare across sets of parameters, we return to the simple synthetic static choice problem presented in equation (12).

Figure 18 plots the resulting optimal allocations for the parameter sets that result from targeting different moments. Across all estimates, LTC expenditure is allocated the majority of wealth. In comparisons across the specifications, several patterns are observable that reflect differences in parameter estimates. We observe that the estimates that exclusively target wealth moments indicate that bequests are more of a luxury good and thus receive a lower share of allocated wealth. In addition, these wealth estimates indicate a very high marginal value of wealth when in the LTC state, which is reflected in figure 18*B* by the slow decline in LTC expenditure as wealth increases. The SSQ estimates indicate high θ_{ADL} and θ_{beq} . The SSQ estimates imply a low marginal value of wealth in these states, which is reflected by the steeper allocation profile of normal consumption in figure 18*C*. Finally, note that the lines in figure 18*A* generally fall between the corresponding lines in the other two panels. Although this does not hold at lower levels of wealth exactly, this broad pattern reflects the trade-off between matching the two sets of moments in the joint estimation. In sum, the SSQ-only, wealth-only, and baseline estimations all result in estimates that indicate a strong desire to spend when in need of help with ADLs.

B. Comparison of Preferences across the Literature

Different preference parameters can induce radically different saving and spending behaviors. To demonstrate the degree of difference and




FIG. 18.—Expenditure fractions across parameter estimates.

to illustrate the large impact of these differences on implied saving behavior, we contrast the implications of the estimated baseline parameters with those in the literature.

In figure 19, we compare the relative strength of these motives induced by the parameters we have estimated to those estimated in De Nardi, French, and Jones (2010; DFJ parameters) and Lockwood (2018), as these are among the papers most closely related to ours. We translate the parameters from other papers to be compatible with our utility function specification, yielding DFJ parameters of $\gamma = 3.81$, $\theta_{ADL} = 0.79$, $\kappa_{ADL} = 0$, $\theta_{beq} = 2,360$, and $\kappa_{beq} = 273$ and Lockwood parameters of $\gamma = 3$, $\theta_{ADL} = 1$, $\kappa_{ADL} = 0$, $\theta_{beq} = 1,460$, and $\kappa_{beq} = 231$. In an analysis similar to that presented in figure 18, in figure 19A we present the optimal bequest allocation (z_2) from the following optimization problem when it is calibrated according to each paper’s baseline estimate of risk aversion and bequest parameters:




FIG. 19.—Expenditure levels across studies: De Nardi, French, and Jones (2010; DFJ) and Lockwood (2018).

$$\begin{aligned} \max_{z_1, z_2} \quad & \frac{(z_1)^{1-\gamma}}{1-\gamma} + \frac{(\theta_{\text{beq}})^\gamma (z_2 + \kappa_{\text{beq}})^{1-\gamma}}{1-\gamma}, \\ \text{subject to } & z_1 + z_2 \leq W, \\ & z_1, z_2 \geq 0, \text{ and } z_2 \geq -\kappa_{\text{beq}}. \end{aligned} \tag{13}$$

We observe that, relative to these other studies, in this paper the preference for leaving a bequest is estimated to be somewhat stronger at low levels of wealth but much weaker at high wealth levels. Whereas these studies' estimates suggest a steadily increasing allocation to bequests even for wealth levels above \$400,000, our estimates suggest that the allocation share is relatively stable for wealth levels above \$100,000. Thus, in our estimation, bequests are considered less of a luxury good, compared to the findings of both De Nardi, French, and Jones (2010) and Lockwood (2018). Furthermore, for high levels of wealth, the bequest share asymptotes to 48% in our baseline, compared to 89% for De Nardi, French, and Jones and 92% for Lockwood. These differences reflect that we estimate a much lower bequest multiplier, and hence a lower marginal value of bequests, at high wealth levels than does either other study. These differences are illustrated just in the simple synthetic allocation problem, but it helps to demonstrate why in our model most of bequeathed wealth is given incidentally rather than as a result of a bequest motive per se.

De Nardi, French, and Jones (2010) show that health risks play an important role in determining late-in-life savings patterns, while Lockwood (2018) shows that highly valued bequests can explain much of late-in-life wealth holding. This difference between the two studies in the preference for leaving a bequest can be large at the lower wealth levels featured in their samples. At higher wealth levels, these differences shrink, documenting the value of the higher-wealth VRI sample in identifying bequest motives. Indeed, De Nardi, French, and Jones are well aware of this, writing "Our sample of singles may not contain enough rich households to reveal strong bequest motives" (62). We can extend De Nardi, French, and Jones's claims that "most people in [their] sample do not have a strong bequest motive" (62) by documenting that a sample with higher-wealth individuals yields a similar lack of individuals with a strong preference to leave a large bequest.

One reason we estimate weaker bequest motives is that our modeling of nonhomothetic utility in the ADL health state reduces the need for a strong bequest motive to match saving patterns for middle-wealth individuals. In figure 19*B*, we present the z_2 allocation for the problem presented in equation (13), with the bequest function replaced by the LTC-state utility function. Lockwood (2018) does not have state-dependent utility in the baseline model and hence would suggest an equal allocation across states. De Nardi, French, and Jones (2010) do allow for a differential marginal

utility multiplier when an individual is unhealthy, but no marginal utility additive shifter. (Additionally, De Nardi, French, and Jones allow for health-state-dependent utility when an individual is sick, but they do not model the LTC state separately.) The estimated state-dependent utility parameter in De Nardi, French, and Jones (2010) is insignificant and assigns almost equal marginal utility to both states, resulting in an almost equal allocation. In this paper, we estimate LTC expenditure to be valued strongly as a necessity, resulting in a much higher allocation to the LTC-state expenditure, not only at low levels of wealth but also at higher wealth levels. Furthermore, estimation of a negative κ_{ADL} has the attractive feature that, eventually, the extra spending when in need of LTC driven by this motive diminishes with wealth.

Analyzing the behavioral implications of parameters from the literature.—We have documented substantial differences in the preferences estimated in the literature and that these differences induce very different savings patterns and imply different motives determining saving behavior. Using both new wealth data on high-wealth individuals and new SSQ measurements, we have estimated preferences that imply that savings are driven to a significant degree by health-related precautionary saving motives.

Figure 20 compares moments of the population wealth profiles induced by these model parameters to those in the data. In addition to Lockwood and DFJ parameter values, we include a set of parameters designed to capture the classic setup, similar to that in Yaari (1965), in which there is no health-state-dependent utility function, zero bequest utility, no government-provided care, and no minimum expenditure when in need of LTC ($\gamma = 3$, $\theta_{ADL} = 1$, $\kappa_{ADL} = 0$, $\theta_{beq} = 0$, $\kappa_{beq} \approx \infty$, $\psi_G = 0$, $\omega_G = 0$, and $\chi = 0$). Since our baseline parameters were in part chosen to match these wealth targets, it is no surprise that they do so well, as discussed above. In terms of induced savings profiles, our estimated parameters look most similar to those of Lockwood, with the De Nardi, French, and Jones (2010) and especially the Yaari (1965) preferences leading to less saving over the life cycle than is observed in the data.²⁶ These differences show that the wealth data can be used to help distinguish across preference parameters; however, it is in the SSQ moments that the different implied motives really stand out.

Figure 21 compares parameter-implied and empirical SSQ means across parameter sets. The results are striking, particularly for SSQs 3a–3c, which ask respondents to make a trade-off between spending on themselves when in need of LTC and leaving a bequest. The parameters estimated by De Nardi, French, and Jones (2010) and Lockwood (2018) generate data that match SSQs 1 and 2 well but do not match SSQ 3, as they imply a much-too-high propensity to allocate toward bequests, given

²⁶ Many papers in the literature, including De Nardi, French, and Jones (2010), have estimates that are derived from less affluent samples. Thus, it is not surprising that the estimated parameters do not match moments for the high-wealth individuals in the VRI.




FIG. 20.—Model fit to wealth moments for alternative parameters: De Nardi, French, and Jones (2010; DFJ); Lockwood (2018); and Yaari = Yaari (1965).

their estimated θ_{beq} and κ_{beq} . In contrast, since the Yaari (1965) parameters place zero value on bequests, they overshoot and would predict much too small an allocation toward bequests, compared to the data. Together, figures 20 and 21 show that models can be consistent with the preference information provided by SSQs without sacrificing fit with the empirical wealth-age distribution. Furthermore, SSQ methods can be used as a powerful parameter-identification device.

Given that these parameters represent significantly different preferences and induce very different saving behaviors in different contingencies, SSQs provide extra information that helps to identify these parameter values. Modeling a nonhomothetic LTC-state-dependent utility function allows us to match wealth moments across the wealth and age distribution as well as traditional models, while also being consistent with survey evidence on stated preferences.

IX. Additional Analyses

A. Sensitivity Analysis

In this section, we provide additional analysis that demonstrates the robustness of the main conclusion that LTC-related utility and risks are a

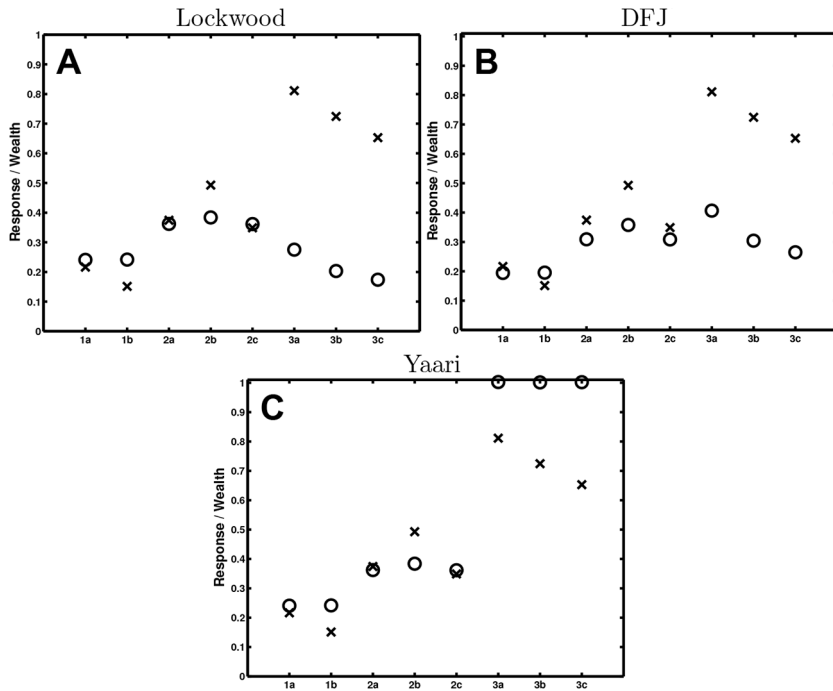


FIG. 21.—Mean SSQ responses in model (circles) and data (crosses) for alternative parameters: A, Lockwood (2018); B, De Nardi, French, and Jones (2010; DFJ); C, Yaari (1965).

significant driver of late-in-life saving behavior. To show this, we perform sensitivity analysis by reestimating the model under the following alternative assumptions: (1) we vary the CRRA parameter (γ); (2) we vary the time discount factor (β); (3) we allow for different curvature exponents in the utility functions for bequests and when alive; (4) we allow for a more severe definition of needing help with ADLs; (5) we eliminate and double out-of-pocket medical expenses; and (6) we add housing wealth to financial wealth and match the model to this measure of total wealth.

Table 6 reports the alternative parameter estimates associated with each sensitivity analysis exercise. In addition, to provide a sense of the behavior that would be induced by the alternative parameters, the table reports expenditure shares associated with the synthetic static allocation problem presented in equation (12). The striking feature across all exercises is the large and negative κ_{ADL} , coupled with $\theta_{ADL} < 1$, which drives a large allocation to ADL expenditure.

TABLE 6
SENSITIVITY ANALYSIS

ALTERNATIVES	PARAMETER ESTIMATES							SYNTHETIC EXPENDITURE SHARES			
								W = \$100K		W = \$200K	
	γ	θ_{ADL}	κ_{ADL}	θ_{beq}	κ_{beq}	ψ_G	ADL	Bequest	ADL	Bequest	
Baseline	5.27	.67	-37.44	1.09	7.83	77.43	.68	.11	.56	.19	
Risk aversion:											
$\gamma = 2$	2.00	.71	-34.49	.45	22.13	58.48	.61	.20	.46	.34	
$\gamma = 8$	8.00	.86	-42.05	1.06	12.82	82.05	.68	.09	.53	.20	
Discount factor:											
$\beta = .98$	4.91	.69	-41.55	.76	12.37	81.55	.69	.12	.54	.24	
$\beta = .99$	4.65	.88	-41.27	1.09	6.92	81.27	.66	.13	.51	.21	
$\beta = .999$	4.90	.85	-39.79	.99	9.10	79.79	.65	.13	.51	.22	
Bequest utility exponent:											
$\gamma_{beq} = 2$	3.01	.61	-36.98	1.23	18.00	69.73	.68	.14	.48	.34	
$\gamma_{beq} = 8$	7.56	.80	-40.59	1.05	12.48	96.74	.69	.08	.55	.17	
ADL-state definition:											
≥ 2 ADLs	5.27	.68	-37.38	1.09	7.79	82.39	.68	.11	.56	.19	
≥ 3 ADLs	5.48	.69	-34.20	.95	8.34	79.20	.65	.14	.53	.22	
Health expenditures:											
None	5.28	.68	-36.80	.99	7.30	81.80	.67	.13	.55	.21	
$2\times$	5.27	.68	-34.35	1.04	7.96	79.35	.66	.13	.54	.20	
Housing wealth:											
15% transaction cost	5.47	.66	-37.27	.89	5.90	82.28	.66	.15	.54	.23	
No transaction cost	6.02	.68	-39.10	.81	9.83	84.10	.67	.14	.53	.24	

1. Different Utility Curvature (γ)

Table 6 presents estimates of all other parameters when setting $\gamma = 2$ or $\gamma = 8$, in contrast to the baseline estimate of $\gamma = 5.27$. When risk aversion is higher, not much changes. With a smaller coefficient of relative risk aversion, there is slightly more need for a stronger bequest utility to match the same wealth data. This appears in the large decrease in θ_{beq} , since κ_{beq} actually increases. This change in preferences nearly doubles bequests in the synthetic allocation problem. Nonetheless, for this wide range of γ , LTC remains by far the dominant expenditure in the simple static allocation problem.

2. Different Time Discount Factor (β)

A higher β , reflecting more-patient individuals, is associated with minor reductions in ADL expenditure and slight increases in bequests. The CRRA parameter estimate is lower than that in the baseline, as increased patience helps match the wealth data with a somewhat weaker precautionary motive. Although there are some changes in bequest and ADL parameters, associated synthetic expenditure shares do not move much. In

total, changes in the discount factor between the baseline $\beta = 0.97$ and $\beta = 0.999$ do not much affect the estimated parameters or the main conclusion that LTC-related preferences and risks are a significant driver of savings over the life cycle.

3. Bequest Utility Exponent (γ_{beq})

In this exercise, we allow for a bequest-specific exponent in the utility function. That is, different from γ , we allow for γ_{beq} such that $v(b) = (\theta_{\text{beq}})^{-\gamma_{\text{beq}}} (b + \kappa_{\text{beq}})^{1-\gamma_{\text{beq}}} / (1 - \gamma_{\text{beq}})$.

Since the parameter is not precisely estimated using wealth data alone and we do not have an SSQ to separately identify a bequest curvature parameter, we explore $\gamma_{\text{beq}} = 2$ and $\gamma_{\text{beq}} = 8$. We choose one value below and one above the ordinary CRRA parameter γ to explore whether the LTC-related saving motive remains strong with a more flexible parameterization. In particular, the lower γ_{beq} makes the bequest utility function less concave and provides more potential for a large bequest motive. *Ceteris paribus*, these changes represent significant differences in the marginal utility of a bequest, relative to the baseline. Nonetheless, when the model is reestimated most parameters are little affected; the ordinary CRRA parameter γ changes the most from its baseline value. Aside from the estimate of γ , qualitatively and quantitatively the other parameter estimates remain similar to baseline values. When $\gamma_{\text{beq}} = 2$, the bequest motive generates more saving, and the model can match the same wealth moments with less precautionary savings, resulting in a lower estimate of $\gamma = 3.01$. At $\gamma_{\text{beq}} = 8$, the precautionary saving motive must be stronger to match the wealth moments, resulting in a larger-than-baseline $\gamma = 7.56$. Expenditure when in need of LTC is still viewed as a strong necessary good with substantial marginal valuation, indicated by a large and negative κ_{ADL} and $\theta_{\text{ADL}} < 1$.

At lower levels of expenditure, for example, $W = 100K$, the synthetic expenditure shares on ADL and bequests remain close to those resulting from the baseline parameters. At higher expenditure levels, for example, \$200K, when $\gamma_{\text{beq}} = 2$, the synthetic expenditure share on bequests is larger than that in the baseline. When γ_{beq} equals 2 or 8, bequests are estimated as more of a luxury good (larger κ_{beq}). The higher $\gamma_{\text{beq}} = 8$ represents a more concave bequest function, which combines with the higher κ_{beq} and similar θ_{beq} to generate less bequests. When the utility function is less concave, at higher wealth levels the lower $\gamma_{\text{beq}} = 2$ more than offsets the higher κ_{beq} to generate an increased desire to leave a bequest.

In sum, whether the bequest function is more or less concave than that in the baseline, whenever an individual needs help with LTC, LTC utility

is always estimated to represent a strong, necessary desire to spend with a high marginal valuation.

4. Definition of the ADL State

When mapping the model to the data, we defined $s = 2$ as meaning “needs and is receiving help with at least one ADL.” In this exercise, we change the definition of ADLs to capture only people who are in need of help with multiple ADLs. First, we classify $s = 2$ if an individual needs and receives help with at least two ADLs and then for at least three ADLs. This has three effects that partially offset each other. First, the ADL state is more severe, with higher persistence in the ADL state, which increases the precautionary saving motive, all other things equal. Second, the more severe ADL state is associated with lower survival rates, reducing expected longevity. Third, individuals are less likely to enter the ADL health state, reducing precautionary saving for any individual and reducing the fraction of people dissaving rapidly in the ADL state. Because of these offsetting effects, the different estimated health-state transition probabilities leave the estimated preference parameters largely unchanged.

5. Different Health Expenditure Risk

Given the difficulty of accurately measuring health expenditure costs and of separating exogenous from endogenous expenditures in the HRS, we explore the effect of large changes in the health expenditure distribution on parameter estimates. We set all medical expenditure shocks to zero and also double all expenditure shocks. In both cases, estimates are substantively unchanged. While these expenses are likely to be more important for the US population, they play a smaller role for people represented by the VRI sample.

6. Housing Wealth and Transaction Costs

In the baseline model, we map the wealth variable a to total financial wealth in the data. We do this because housing is a complicated asset that is in some ways like financial wealth and in some ways more like a durable consumption good. Housing is like financial wealth in that it can be sold and transformed into liquid wealth, albeit in a market subject to frictions. Housing is unlike financial wealth in that it is a physical good that delivers housing services. In this exercise, we map a to the sum of housing and financial wealth. In addition to using the full value of the house, to account for costs associated with the sale of a house, we also do this by subtracting 15% of the house price. Either way, this results in only modest changes in the estimated parameters and associated expenditure shares.

B. External Validity Using HRS Data

Our final exercise is to use the estimated baseline parameters to compare savings patterns in the data and the model, using the US-representative HRS population. In this section, we provide a summary of this analysis, with details provided in the online appendix. First, we estimate health transitions, survival probabilities, health costs, and income profiles, using the HRS sample, as opposed to the VRI-eligible HRS subsample. Then, we sample single households from the HRS and record age, health status, and permanent income quintile in the year in which they enter the HRS. We generate a sequence of health transitions and cost shocks for each household and simulate their saving and expenditure patterns, using the optimal policies associated with our estimated baseline model. Finally, we compare percentiles of the simulated wealth distribution with those measured by the HRS in the subsequent waves. When possible, we do this analysis separately for each cohort, given the different sampling procedures and aggregate environment associated with each cohort.

Overall, our model with our baseline parameter estimates performs well in the HRS sample and matches the 25th, 50th, and 75th percentiles of the wealth distribution for the cohorts that we consider, with a few notable exceptions. For the War Baby and Early Baby Boomer cohorts, our model matches closely the wealth distribution by age. The notable exception is that the model overpredicts wealth in the 2010 wave, possibly because we do not model the financial crisis. When using the Asset and Health Dynamics of the Oldest Old (AHEAD) sample, the model again matches the data well, although it slightly overpredicts the spend-down of wealth in the median wealth percentile later in life. This analysis shows that even though our model is estimated using data from the VRI that is not representative of the US population, our estimated preferences and assumptions about incomplete markets result in a model that is informative about the motives for saving in the more general US population.

X. Conclusion

Individuals face multiple risks and have multiple objectives late in life. The risks include longevity, health, and the need for care. The objectives include sustaining consumption in potentially different health states and providing a bequest. With the decline of defined-benefit pensions, with increased longevity, with high and variable cost of care, and absent or with highly imperfect insurance markets for late-in-life risks, older households are increasingly responsible for managing financial assets late in life to address these multiple objectives across states of the world.

In this paper, we present novel preference elicitation methods and data to answer questions about late-in-life saving and spending. The modeling, measurement, and data were designed in tandem.

- We build an incomplete-markets model of individuals, who save precautionarily when faced with health risks, the potential need for LTC, and an uncertain life span and who value consuming, leaving a bequest, and receiving LTC if they need it. Expenditures when in need of LTC can be valued differently than ordinary consumption, depending on estimates of a nonhomothetic health-state-dependent utility function, and individuals choose on the margin the amount to spend when in need of LTC.
- We develop SSQs that use a novel application of stated-preference methodology. These SSQs are designed to elicit responses that identify the preference parameters of the model.
- We develop a novel data infrastructure—the VRI—that combines information on income, assets, and demographics with the SSQ responses on a large sample of older Americans with sufficient financial assets to make operational the financial trade-offs implied by the model.

We estimate the parameters of the model, using moments from the wealth distribution alone, SSQs alone, and both wealth and SSQs. The point estimates based on the traditional approach of using the wealth distribution and those based on the novel approach of using the SSQs are generally similar. The estimates with the wealth data alone are, however, imprecise. Because multiple motives generate wealth accumulation and wealth is fungible across outcomes, wealth data alone do not strongly identify motives. In contrast, the SSQs provide relatively sharp identification of relevant parameters. This precise identification of parameters allows us to make inferences that are difficult to support with behavioral data alone.

The findings in the paper point toward the risk of needing LTC when old as a substantial motive for accumulating assets. The marginal utility of expenditures when in need of LTC is larger than that of bequests. Because of the strength of the estimated health-state-dependent utility function, the precautionary saving motives associated with LTC contribute significantly to late-in-life saving behavior, strongly affecting wealth accumulation—tending to increase savings both across the wealth distribution and over the life cycle—for a large fraction of the US population. For a typical wealth holder in our sample, peak assets in retirement are about twice what they would be absent a difference in the utility of spending between the healthy and need-for-care states. Health-state-dependent utility also helps explain the postretirement pattern of assets. It is a

motive for continued asset accumulation in retirement. LTC-associated expenses also lead to a rapid drawdown of assets late in life and therefore generate lower bequests, on average, than the continued buildup of assets would otherwise suggest.

The model- and SSQ-based estimates help distinguish the motives for leaving bequests—something difficult to do with the behavioral data alone because many bequests will be incidental, that is, occur in the event where an earlier-than-expected death or a less-costly-than-expected old age results in a larger-than-expected bequest. Bequests in the baseline model resulting from longevity risk, precautionary saving for LTC risk, and a warm-glow bequest motive are, on average, more than twice the incidental bequest owing to longevity uncertainty alone. Overall, we show that the desire to self-insure against LTC risk explains a substantial fraction of the wealth holding of older Americans.

Finally, if LTC and longevity risks are substantial and if people have a desire to spend when in need of care and to leave a bequest, why do people not buy large amounts of LTCI and annuities? There is low take-up of LTCI and very low take-up of privately purchased annuities in the United States. This low take-up is a puzzle. Demand for idealized LTCI in this model would be high and demand for idealized annuities even higher. The model in this paper addresses wealth accumulation in the realistic case that these idealized state-contingent assets are absent from the market and therefore from households' portfolios. In other work, we address the LTCI puzzle. Ameriks et al. (2018c) provide detailed analysis of demand for LTCI, showing that a substantial part of low ownership is explained by state-contingent preferences and incomplete markets; another substantial part is explained by differences between the real-world insurance products and the idealized insurance products in the model, while some difference remains attributable to model misspecification and measurement error. Absent well-functioning markets for state-contingent assets, the focus of this paper on wealth accumulation to meet desires related to LTC and bequests is highly relevant for understanding behavior.

Appendix A

Estimated Inputs for Model

A1. *Income*

We estimate a deterministic income process from the cross-sectional income distribution. Income is defined as the sum of labor income, publicly and privately provided pensions, and disability income, as measured in VRI Survey 1. The income processes are estimated to be a function of a constant, age, age squared, sex, and the interaction of sex and age. To ensure that income is positive in all periods, we estimate a quantile regression of log income on these variables.

Because we allow for five income profiles, the quantile regression is estimated for the 10th, 30th, 50th, 70th, and 90th percentiles of the income distribution. We calibrate our income processes to the resulting estimates and group individuals into income profile quintiles.

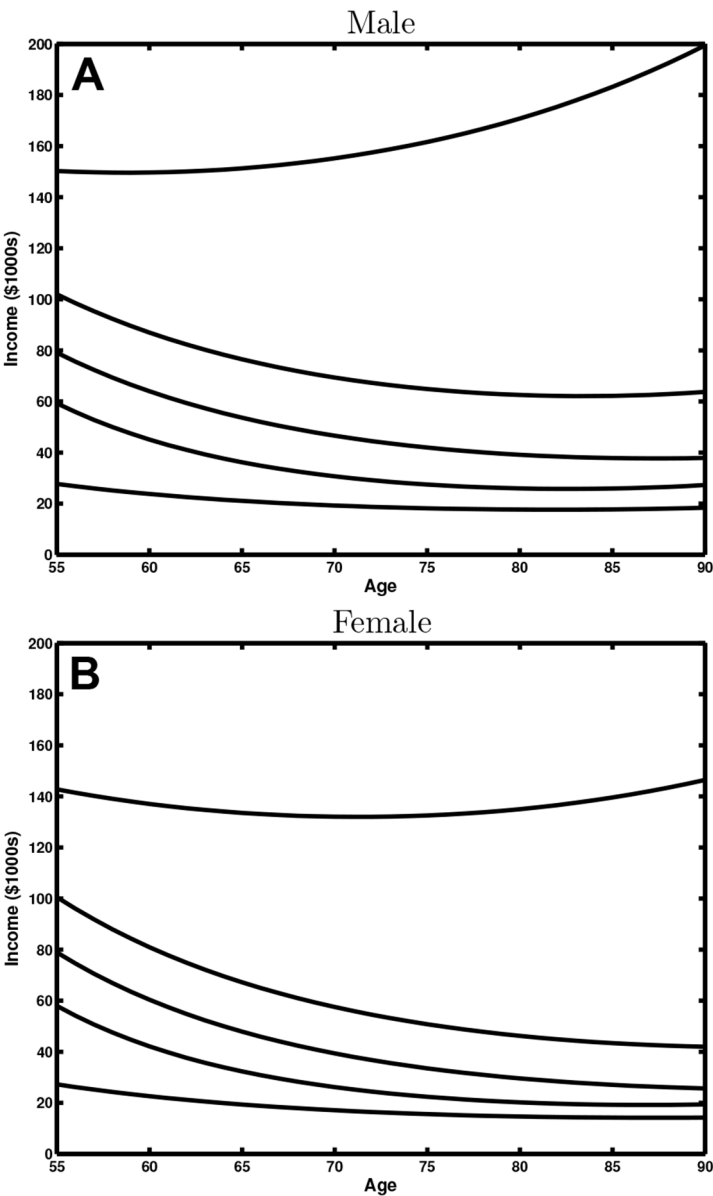


FIG. A1.—Income profile quintiles for males (A) and females (B).

A2. Health

A2.1. Health-State Transition Matrix

Using appropriate health-state definitions, we estimate a sequence of health-transition matrices conditional on a vector $x_{i,t}$ that includes individual i 's age, t , and sex, g . The HRS records only 2-year health-state transitions, which we use to identify the one-year transition probabilities in a manner similar to De Nardi, French, and Jones (2010). To do this, we write the 2-year transition probabilities as

$$\begin{aligned}\Pr(s_{t+2} = j | s_t = i) &= \sum_{k=0}^3 \Pr(s_{t+2} = j | s_{t+1} = k) \Pr(s_{t+1} = k | s_t = i) \\ &= \sum_{k=0}^3 \pi_{kj,t+1} \pi_{ik,t},\end{aligned}$$

where

$$\pi_{ik,t} = \frac{\phi_{ik,t}}{\sum_{m=0}^3 \phi_{im,t}} \text{ and } \phi_{ik,t} = \exp(x_{i,t} \beta_k).$$

We then estimate β_k , using a maximum likelihood estimator, and use these estimates to construct the corresponding cells in the health-transition matrices. Health transitions are estimated using HRS waves 2–10. The questions necessary to make the health-state assignment are not available in the 1992 survey, so we exclude this wave from the health-transition estimates.

Figures A2 and A3 display the estimated health-state transition probabilities ($\pi_g(s'|t, s)$). Section V.A.2 describes the estimation methodology. An additional consideration is how to define the “needs-LTC” health status. There are three measures in the HRS that could potentially be used. The first is nursing home stay, the second is needs help with the ADLs, and the third is receives help with the ADLs.

Nursing home stay (more than 120 nights in a nursing home before the current interview or currently in a nursing home at time of interview) is what De Nardi, French, and Jones (2010) used. Given that we allow people in the model to choose their type of care, we want a less restrictive definition for $s = 2$. The ADL questions in the RAND version of the HRS list many ADLs and asks whether the respondent has difficulty completing those tasks without help. In some sense, these questions provide the broadest possible definition of the ADL state, since many people could report having difficulty but would still be able to live without receiving help. We choose to implement the intermediate measure: we categorize and individual as needing help with LTC if they have difficulty with at least one ADL and they also receive outside help completing the ADL task. Since we are not using stays in a nursing home to represent our health state $s = 2$, not modeling spending when in need of care as an out-of-pocket expenditure shock, and not using a US-representative sample, our model-generated health and spending patterns will be different from those in Hurd, Michaud, and Rohwedder (2017) and Friedberg et al. (2014).




FIG. A2.—Male health-state transition profiles: healthy (A), sick (B), LTC (C).

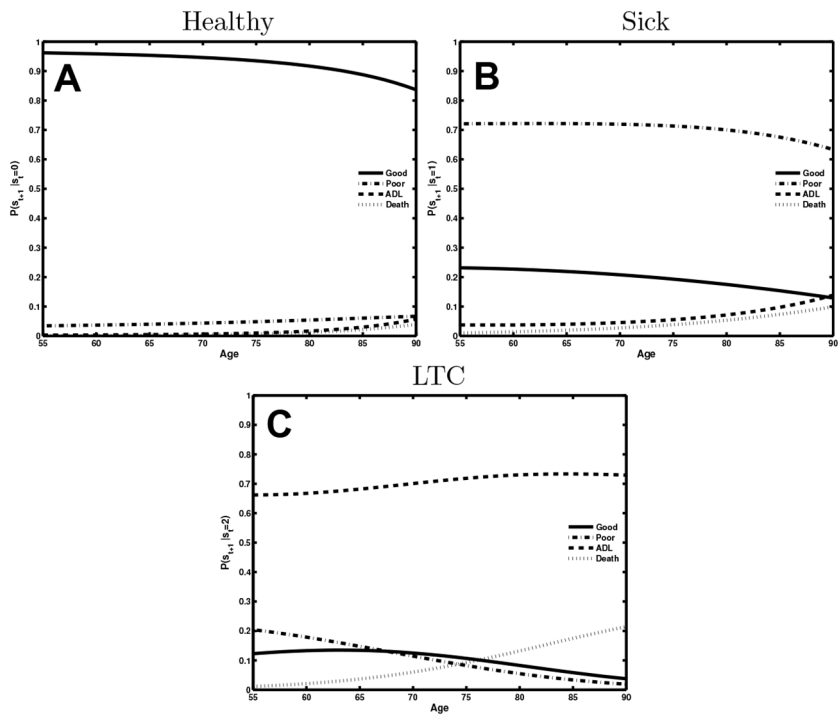


FIG. A3.—Female health-state transition profiles: healthy (A), sick (B), LTC (C).

A2.2. Health Cost

To estimate the mean of the health cost distribution, $\mu(t, g, s)$, we run an OLS regression of annualized log out-of-pocket medical expenditures (variable `r10oopmd` in the RAND HRS) on age, sex, health state, and interaction terms. Log medical expenses are modeled as a linear function of a quartic in age, prior health status, sex, and health interacted with age. Using the residuals from this first regression, we regress the squared residuals on the same set of state variables as in the first regression to find the conditional variance of medical expenses, $\sigma^2(t, g, s)$. The resulting mean and variance estimates parameterize the lognormal distribution for medical expenses conditional on age, sex, and health status. We then discretize the error term, using the Tauchen method to generate the medical expense process.

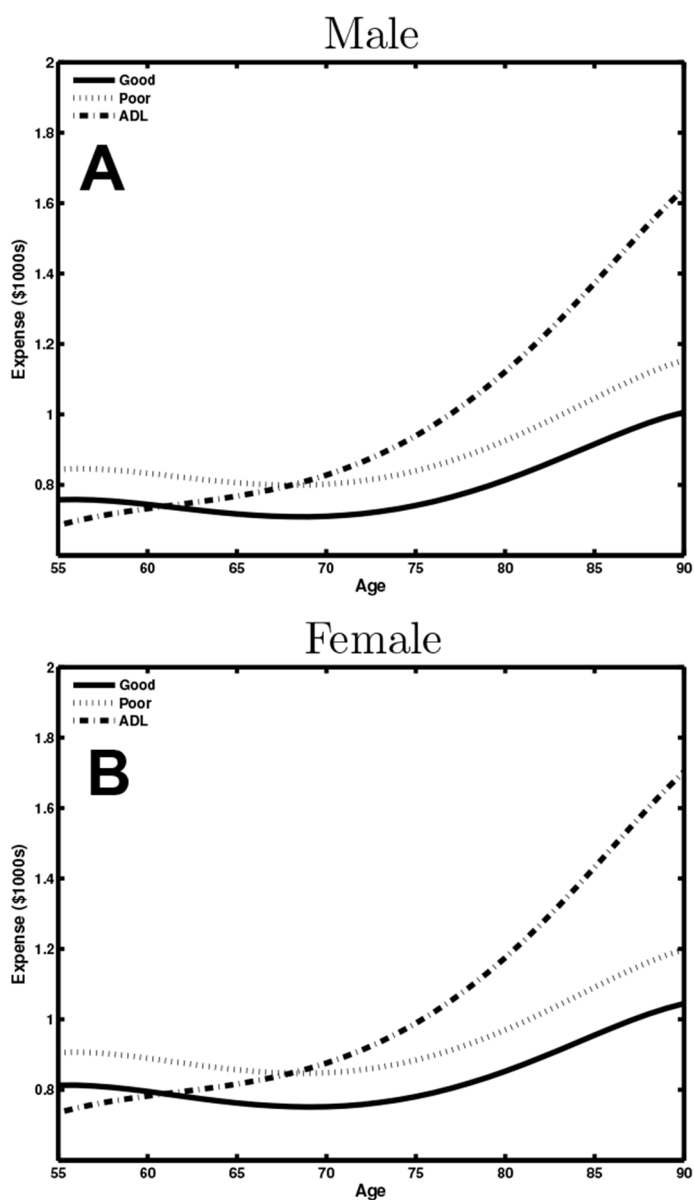


FIG. A4.— Median health cost profile.

Figure A4 plots the median health cost shock over the life cycle by men and women of different health statuses. Men in poor health spend around \$100 more per year for health costs than healthy men. Later in life, men in need of LTC spend about \$600 more than healthy men. Saving is also affected by the risk of high costs, so in table A1 we present examples of the health cost distribution across sex, age, and health states. Overall, these health costs are much smaller than LTC expenditures and thus contribute little to the overall saving motive. The ideal health cost measure would involve some fine-tuning, for example, removing some mandatory component of nursing home costs and adding end-of-life expenses. Since in our sensitivity analysis we see little effect of setting these costs to zero or doubling these costs, we do not dive more deeply into health cost spending subcategories. See French, Jones, and McCauley (2017) for detailed analysis of health-related spending in the HRS and the MCBS. Since we model spending when in need of LTC, instead of spending on LTC, there is no direct mapping from output in our model to health spending in these surveys.

TABLE A1
MEAN AND STANDARD DEVIATION OF THE HEALTH EXPENSE LOGNORMAL
DISTRIBUTIONS (\$000s)

HEALTH STATE	AGE									
	55	60	65	70	75	80	85	90	95	100
Female Health Expenses										
Healthy ($s = 0$), mean	2.4	2.0	1.8	1.7	1.9	2.4	3.1	3.9	4.2	3.1
SD	6.5	4.7	3.7	3.6	4.2	6.1	9.6	14.2	16.2	10.9
Sick ($s = 1$), mean	3.4	2.9	2.6	2.4	2.5	3.1	4.0	4.9	5.1	3.7
SD	12.4	8.7	6.6	6.1	6.8	9.3	13.9	19.7	21.4	13.8
Needs LTC ($s = 2$), mean	3.5	3.1	2.8	2.9	3.3	4.2	5.7	7.5	8.2	6.3
SD	16.6	12.1	9.5	8.9	10.4	14.4	22.1	32.1	36.0	23.8
Male Health Expenses										
Healthy ($s = 0$), mean	2.8	2.3	1.9	1.7	1.8	2.1	2.6	3.1	3.1	2.2
SD	10.1	6.5	4.6	3.8	4.0	5.0	6.9	9.0	9.1	5.3
Sick ($s = 1$), mean	4.1	3.2	2.6	2.4	2.4	2.7	3.3	3.9	3.8	2.6
SD	19.1	11.9	8.0	6.5	6.4	7.7	10.1	12.6	12.1	6.8
Needs LTC ($s = 2$), mean	4.2	3.5	3.0	2.9	3.1	3.7	4.8	5.9	6.2	4.5
SD	25.5	16.4	11.5	9.5	9.7	11.9	16.1	20.6	20.3	11.7

References

Ameriks, John, Joseph Briggs, Andrew Caplin, Matthew D. Shapiro, and Christopher Tonetti. 2018a. “Long-Term Care Model.” Vanguard Res. Initiative Tech. Report. <http://ebp-projects.isr.umich.edu/VRI/papers/VRI-TechReport-LTC-Model.pdf>.

———. 2018b. “Long-Term Care Strategic Survey Questions.” Vanguard Res. Initiative Tech. Report. <http://ebp-projects.isr.umich.edu/VRI/papers/VRI-TechReport-LTC-SSQs.pdf>.

———. 2018c. “The Long-Term-Care Insurance Puzzle: Modeling and Measurement.” Working Paper no. 22726 (June), NBER, Cambridge, MA.

Ameriks, John, Andrew Caplin, Steven Laufer, and Stijn Van Nieuwerburgh. 2011. “The Joy of Giving or Assisted Living? Using Strategic Surveys to Separate Public Care Aversion from Bequest Motives.” *J. Finance* 66 (2): 519–61.

- Ameriks, John, Andrew Caplin, Minjoon Lee, Matthew D. Shapiro, and Christopher Tonetti. 2015. "The Wealth of Wealthholders." Working Paper no. 20972 (February), NBER, Cambridge, MA.
- Arrow, Kenneth J. 1974. "Optimal Insurance and Generalized Deductibles." *Scandinavian Actuarial J.* 1974:1–42.
- Barsky, Robert B., F. Thomas Juster, Miles S. Kimball, and Matthew D. Shapiro. 1997. "Preference Parameters and Behavioral Heterogeneity: An Experimental Approach in the Health and Retirement Study." *Q.J.E.* 112 (2): 537–79.
- Blass, Asher A., Saul Lach, and Charles F. Manski. 2010. "Using Elicited Choice Probabilities to Estimate Random Utility Models: Preferences for Electricity Reliability." *Internat. Econ. Rev.* 51 (2): 421–40.
- Braun, R. Anton, Karen A. Kopecky, and Tatyana Koreschkova. 2017. "Old, Sick, Alone, and Poor: A Welfare Analysis of Old-Age Social Insurance Programmes." *Rev. Econ. Studies* 84 (2): 580–612.
- Brown, Jeffrey R., and Amy Finkelstein. 2008. "The Interaction of Public and Private Insurance: Medicaid and the Long-Term Care Insurance Market." *A.E.R.* 98 (3): 1083–102.
- Brown, Jeffrey R., Gopi Shah Goda, and Kathleen McGarry. 2016. "Heterogeneity in State-Dependent Utility: Evidence from Strategic Surveys." *Econ. Inquiry* 54 (2): 847–61.
- Chetty, Raj, Michael Stepner, Sarah Abraham, et al. 2016. "The Association between Income and Life Expectancy in the United States, 2001–2014." *JAMA* 315 (16): 1750–66.
- De Nardi, Mariacristina. 2004. "Wealth Inequality and Intergenerational Links." *Rev. Econ. Studies* 71 (3): 743–68.
- De Nardi, Mariacristina, Eric French, and John B. Jones. 2010. "Differential Mortality, Uncertain Medical Expenses, and the Saving of Elderly Singles." *J.P.E.* 118 (1): 39–75.
- . 2016a. "Medicaid Insurance in Old Age." *A.E.R.* 106 (11): 3480–520.
- . 2016b. "Savings After Retirement: A Survey." *Ann. Rev. Econ.* 8:177–204.
- Dynan, Karen E., Jonathan Skinner, and Stephen P. Zeldes. 2002. "The Importance of Bequests and Life-Cycle Savings in Capital Accumulation: A New Answer." *A.E.R.* 92 (2): 274–78.
- Evans, William N., and W. Kip Viscusi. 1991. "Estimation of State-Dependent Utility Functions Using Survey Data." *Rev. Econ. and Statis.* 73 (1): 94–104.
- Fella, Giulio. 2014. "A Generalized Endogenous Grid Method for Non-smooth and Non-concave Problems." *Rev. Econ. Dynamics* 17 (2): 329–44.
- Finkelstein, Amy, Erzo F. P. Luttmer, and Matthew J. Notowidigdo. 2009. "Approaches to Estimating the Health State Dependence of the Utility Function." *A.E.R.* 99 (2): 116–21.
- . 2013. "What Good Is Wealth without Health? The Effect of Health on the Marginal Utility of Consumption." *J. European Econ. Assoc.* 11 (s1): 221–58.
- French, Eric, and John B. Jones. 2011. "The Effects of Health Insurance and Self Insurance on Retirement Behavior." *Econometrica* 79 (3): 693–732.
- French, Eric, John B. Jones, and Jeremy McCauley. 2017. "The Accuracy of Economic Measurement in the Health and Retirement Study." *Forum Health Econ. and Policy* 20(2). <https://doi.org/10.1515/fhep-2017-0001>.
- Friedberg, L., W. Hou, W. Sun, A. Webb, and Z. Li. 2014. "New Evidence on the Risk of Requiring Long-Term Care." Working Paper no. 2014–12, Boston Coll. Center Retirement Res.
- Genworth. 2016. "Genworth 2016 Cost of Care Study." Discussion paper, Genworth, Richmond, VA.

- Gourinchas, Pierre-Olivier, and Jonathan A. Parker. 2002. "Consumption over the Life Cycle." *Econometrica* 70 (1): 47–89.
- Gustman, Alan L., and Thomas L. Steinmeier. 1986. "A Structural Retirement Model." *Econometrica* 54 (3): 555–84.
- Hong, Jay H., Josep Pijoan-Mas, and José Victor Ríos-Rull. 2015. "Health Heterogeneity and Preferences." Working paper.
- Hubbard, R. Glenn, Jonathan Skinner, and Stephen P. Zeldes. 1994. "Expanding the Life-Cycle Model: Precautionary Saving and Public Policy." *A.E.R.* 84 (2): 174–79.
- . 1995. "Precautionary Saving and Social Insurance." *J.P.E.* 103 (2): 360–99.
- Hurd, Michael D. 1989. "Mortality Risk and Bequests." *Econometrica* 57 (4): 779–813.
- Hurd, Michael D., Pierre-Carl Michaud, and Susann Rohwedder. 2017. "Distribution of Lifetime Nursing Home Use and of Out-of-Pocket Spending." *Proc. Nat. Acad. Sci. USA* 114 (37): 9838–42.
- Kimball, Miles S., Claudia R. Sahm, and Matthew D. Shapiro. 2008. "Imputing Risk Tolerance from Survey Responses." *J. American Statist. Assoc.* 103 (483): 1028–38.
- Koijen, Ralph S. J., Stijn Van Nieuwerburgh, and Motohiro Yogo. 2016. "Health and Mortality Delta: Assessing the Welfare Cost of Household Insurance Choice." *J. Finance* 71 (2): 957–1010.
- Kopczuk, Wojciech, and Joseph Lupton. 2007. "To Leave or Not to Leave: The Distribution of Bequest Motive." *Rev. Econ. Studies* 74 (1): 207–35.
- Kopecky, Karen A., and Tatyana Koreshkova. 2014. "The Impact of Medical and Nursing Home Expenses on Savings." *American Econ. J.: Macroeconomics* 6 (3): 29–72.
- Kotlikoff, Laurence J., and Lawrence H. Summers. 1981. "The Role of Intergeneration Transfers in Aggregate Capital Accumulation." *J.P.E.* 89 (4): 706–32.
- Laibson, David, Andrea Repetto, and Jeremy Tobacman. 2007. "Estimating Discount Functions with Consumption Choices over the Lifecycle." Working Paper no. 13314 (August), NBER, Cambridge, MA.
- Lillard, Lee A., and Yoram Weiss. 1997. "Uncertain Health and Survival: Effects on End-of-Life Consumption." *J. Bus. and Econ. Statist.* 15 (2): 254–68.
- Lockwood, Lee M. 2018. "Incidental Bequests and the Choice to Self-Insure Late-Life Risks." *A.E.R.* 108 (9): 2513–50.
- Love, David A., Michael G. Palumbo, and Paul A. Smith. 2009. "The Trajectory of Wealth in Retirement." *J. Public Econ.* 93 (1–2): 191–208.
- Palumbo, Michael G. 1999. "Uncertain Medical Expenses and Precautionary Saving near the End of the Life Cycle." *Rev. Econ. Studies* 66 (2): 395–421.
- Poterba, James M., Steven F. Venti, and David A. Wise. 2013. "Health, Education, and the Post-retirement Evolution of Household Assets." Working Paper no. 18695 (January), NBER, Cambridge, MA.
- Scholz, John Karl, Ananth Seshadri, and Surachai Khitatrakun. 2006. "Are Americans Saving 'Optimally' for Retirement?" *J.P.E.* 114 (4): 607–43.
- Soto, Mauricio, Rudolph G. Penner, and Karen E. Smith. 2009. "How Seniors Change Their Asset Holdings during Retirement." Working paper. <https://ssrn.com/abstract=1553268>.
- van der Klaauw, Wilbert. 2012. "On the Use of Expectations Data in Estimating Structural Dynamic Choice Models." *J. Labor Econ.* 30 (3): 521–54.
- van der Klaauw, Wilbert., and Kenneth I. Wolpin. 2008. "Social Security and the Retirement and Savings Behavior of Low-Income Households." *J. Econometrics* 145 (1–2): 21–42.

- Viscusi, W. Kip, and William N. Evans. 1990. "Utility Functions That Depend on Health Status: Estimates and Economic Implications." *A.E.R.* 80 (3): 353–74.
- Yaari, Menahem E. 1965. "Uncertain Lifetime, Life Insurance, and the Theory of the Consumer." *Rev. Econ. Studies* 32 (2): 137–50.